KU LEUVEN

Evasive and efficient distributed adversarial attacks using PSO

Intermediate presentation

Sander Prenen

Thesis supervisors: Prof. dr. ir. W. Joosen, Dr. ir. D. Preuveneers

Mentors: I. Tsingenopoulos, V. Rimmer

0 Outline

1. Background & related work
2. Research topic
3. Progress
4. Evaluation plan

1 Adversarial Attacks

Imperceptibly small perturbations to a correctly classified input image, so that it is no longer classified correctly. [1]

- White box attacks
- Black box attacks
  - Subset of white box attacks
  - More relevant in security use-cases
    - Bypassing malware detection [2]
    - Bypassing face recognition [3]
    - Altering traffic signs [4]

1 Related work

- Boundary attack (BA)

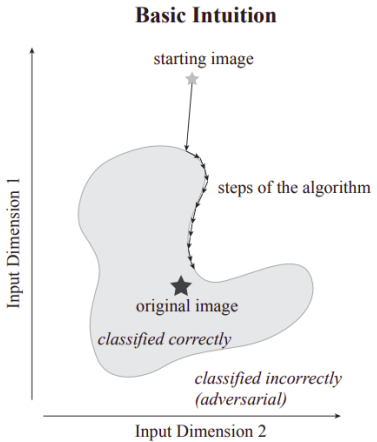

Figure: Boundary attack [5]

1 Related work

- HopSkipJump attack (HSJA)

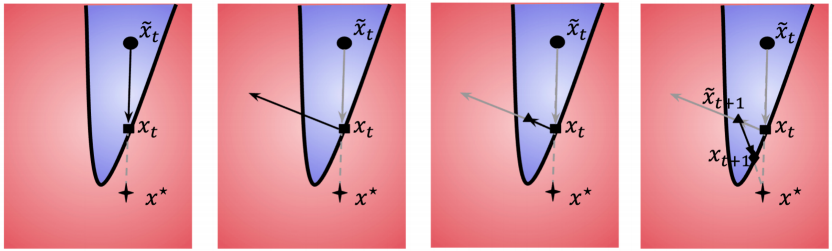

Figure: HopSkipJump attack [6]

1 Adversarial Defenses

- Adversarial training
- Gradient hiding
- Denoising

1 Related work

- Stateful detection

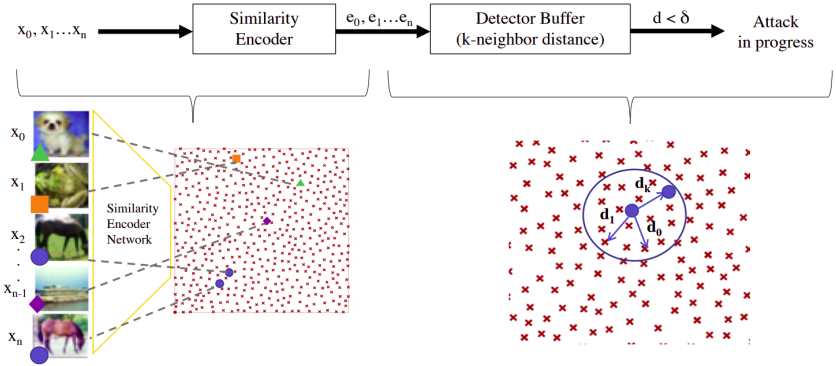

Figure: Stateful detection [7]

1 Related work

- Stateful detection
  - Assumption: attack done by one user/account/IP
  - User can be uniquely identified
  - No cooperation between users

1 Particle Swarm Optimization (PSO)

- Evolutionary algorithm
- Optimization framework
- Inspired by flocking of birds
  - Each particle has a position and corresponding fitness
  - Move based on personal best position, group best position and inertia

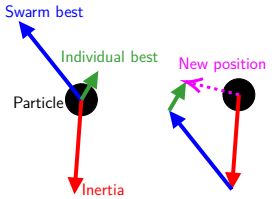

Figure: PSO logic, inspired by [8]

2 Research topic

- Evasive and Efficient Attacks
  - Evade stateful defense
  - By being efficient (less queries)
  - By distribution
- Distribution
  - Centralize the algorithm
  - Distribute the submission of queries
  - Distribute points of attack

2 Research gap

- Distribution
  - Dual goal
    - Evade detection
    - Improve existing attacks using PSO
- Existing work
  - Uses PSO as algorithm in itself
  - Does not evaluate against stateful detection
  - Uses confidence scores [9, 10]

2 Possible research questions

What are the advantages of distributing an adversarial attack?

How can attackers cooperate in order to evade a stateful detection mechanism?

What are the (dis)advantages of using PSO in relation to vanilla adversarial attacks?

3 Threat model

- Decision based attack
- Targeted attack
  - Both are more relevant in real scenarios
- Stateful detection mechanism
- Goals: evade detection & craft best adversarial example

3 Why PSO?

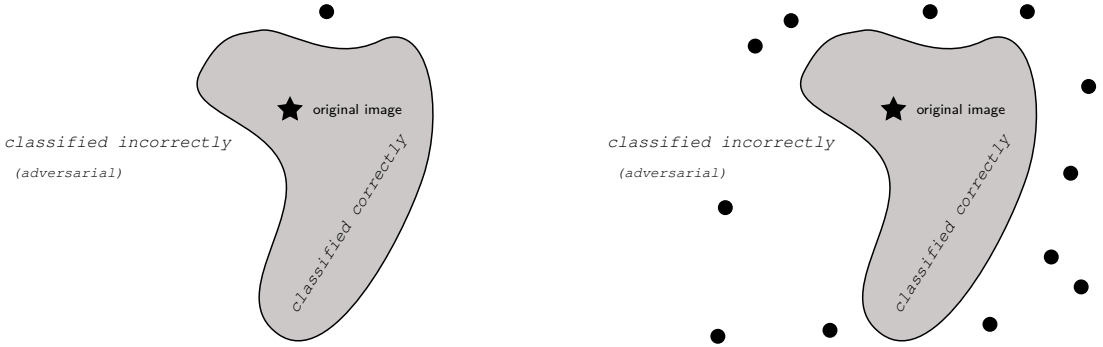

Figure: Advantage of PSO, inspired by [5]

3 Progress

- Working PSO algorithm based on boundary attack (PSO-BA)

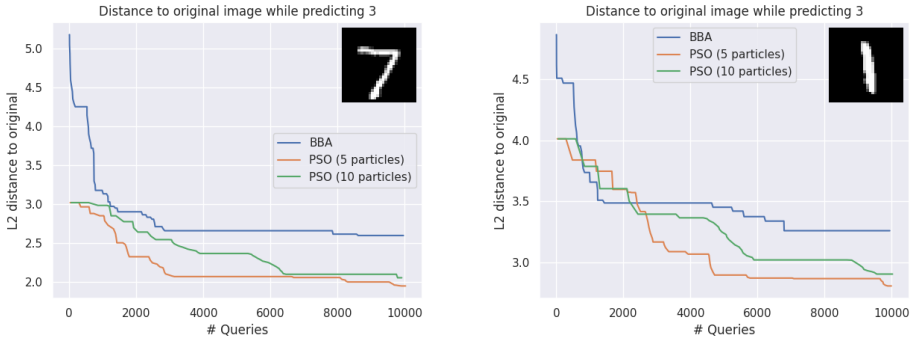

Figure: Comparison BA and PSO-BA

3 Why PSO?

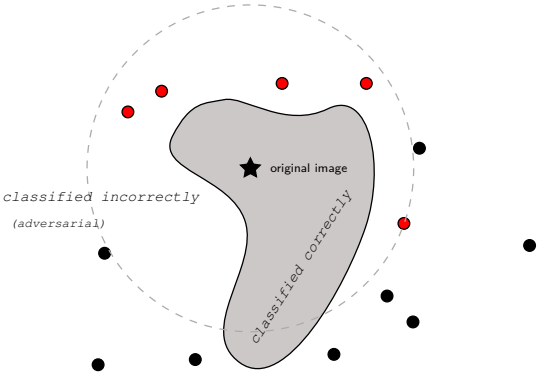

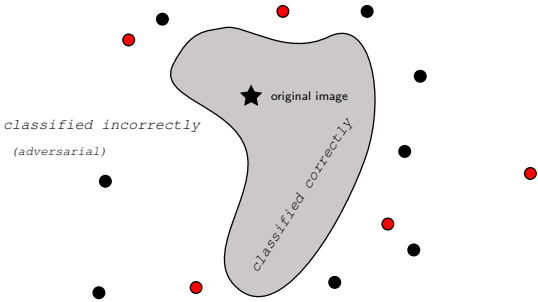

Figure: Advantage of PSO, inspired by [5]

3 Progress

- Working PSO algorithm based on boundary attack (PSO-BA)
- Performed experiments to show that PSO is a viable candidate

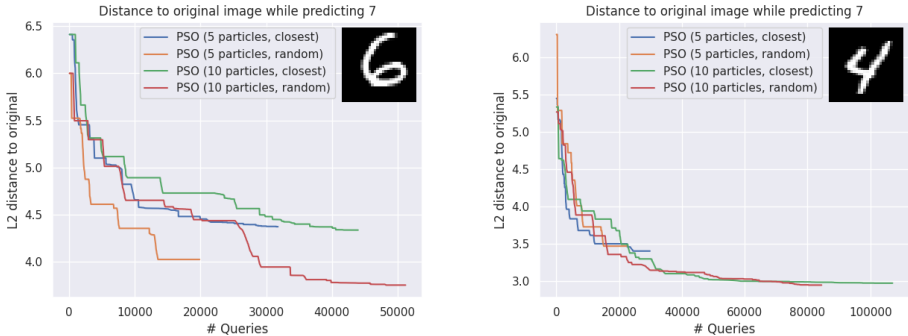

Figure: Comparison random versus closest initialization

3 Progress

- Working PSO algorithm based on boundary attack (PSO-BA)
- Performed experiments to show that PSO is a viable candidate
- Compare detections PSO-BA and BA

	Avg. L_2-distance	Avg. # Detections	Avg. # Queries
BBA	2.9868	148	25010
PSO-BBA	2.8841	79	24721
D-PSO-BBA	2.8841	51	24721

4 Next steps

- Improve the existing PSO-BA algorithm
- Use different methods of distribution
  - Round robin
  - Distance based
  - Other
- Implement a new algorithm based on HSJA and PSO

4 Evaluation plan

- Metric: number of detections and L_2-distance
- Different distribution schemes
- Tuning the hyperparameters
- Applying the algorithm on different datasets
  - CIFAR
  - ImageNet
- Performing more experiments to confirm the results

Questions?

5 References I

- Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus.
  Intriguing properties of neural networks, 2014.
- Octavian Suciu, Scott E. Coull, and Jeffrey Johns.
  Exploring adversarial examples in malware detection, 2019.
- Fatemeh Vakhshiteh, Ahmad Nickabadi, and Raghavendra Ramachandra.
  Adversarial attacks against face recognition: A comprehensive study, 2021.
- Abhiram Gnanasambandam, Alex M. Sherman, and Stanley H. Chan.
  Optical adversarial attack, 2021.

5 References II

Wieland Brendel, Jonas Rauber, and Matthias Bethge.
Decision-based adversarial attacks: Reliable attacks against black-box machine learning models, 2018.

Jianbo Chen, Michael I. Jordan, and Martin J. Wainwright.
Hopskipjumpattack: A query-efficient decision-based attack, 2020.

Steven Chen, Nicholas Carlini, and David Wagner.
Stateful detection of black-box adversarial attacks, 2019.

Nathan Rooy.
Particle swarm optimization from scratch with python.
https://nathanrooy.github.io/posts/2016-08-17/simple-particle-swarm-optimization-with-python/, 08 2016.

5 References III

Rayan Mosli, Matthew Wright, Bo Yuan, and Yin Pan.
They might not be giants crafting black-box adversarial examples using particle swarm optimization.
In Liqun Chen, Ninghui Li, Kaitai Liang, and Steve Schneider, editors, *Computer Security – ESORICS 2020*, pages 439–459, Cham, 2020. Springer International Publishing.

Naufal Suryanto, Hyoeun Kang, Yongsu Kim, Youngyeo Yun, Harashta Tatimma Larasati, and Howon Kim.
A distributed black-box adversarial attack based on multi-group particle swarm optimization.
Sensors, 20(24), 2020.