

Distributed Adversarial Attacks

Sander Prenen

Thesis voorgedragen tot het behalen
van de graad van Master of Science
in de ingenieurswetenschappen:
computerwetenschappen, hoofdoptie
Artificiële intelligentie

Promotoren:

Prof. dr. ir. W. Joosen
Dr. ir. D. Preuveneers

Assessoren:

Ir. W. Eetveel
W. Eetrest

Begeleiders:

V. Rimmer
I. Tsingenopoulos

© Copyright KU Leuven

Without written permission of the thesis supervisors and the author it is forbidden to reproduce or adapt in any form or by any means any part of this publication. Requests for obtaining the right to reproduce or utilize parts of this publication should be addressed to the Departement Computerwetenschappen, Celestijnenlaan 200A bus 2402, B-3001 Heverlee, +32-16-327700 or by email info@cs.kuleuven.be.

A written permission of the thesis supervisors is also required to use the methods, products, schematics and programmes described in this work for industrial or commercial use, and for submitting this publication in scientific contests.

Zonder voorafgaande schriftelijke toestemming van zowel de promotoren als de auteur is overnemen, kopiëren, gebruiken of realiseren van deze uitgave of gedeelten ervan verboden. Voor aanvragen tot of informatie i.v.m. het overnemen en/of gebruik en/of realisatie van gedeelten uit deze publicatie, wend u tot het Departement Computerwetenschappen, Celestijnenlaan 200A bus 2402, B-3001 Heverlee, +32-16-327700 of via e-mail info@cs.kuleuven.be.

Voorafgaande schriftelijke toestemming van de promotoren is eveneens vereist voor het aanwenden van de in deze masterproef beschreven (originele) methoden, producten, schakelingen en programma's voor industrieel of commercieel nut en voor de inzending van deze publicatie ter deelname aan wetenschappelijke prijzen of wedstrijden.

Preface

Sander Prenen

Contents

Preface	i
Contents	ii
List of Figures	iv
List of Tables	v
List of Abbreviations	vi
1 Introduction	1
2 Background	3
2.1 Neural networks	3
2.2 Adversarial attacks	6
2.3 Particle swarm optimization	9
3 Related work	15
3.1 Boundary attack	15
3.2 HopSkipJumpAttack	17
3.3 Stateful defense	19
3.4 PSO and distributed attacks	20
4 Approach	23
4.1 Distribution	23
4.2 Optimization	24
4.3 Approach	24
4.4 Threat model	24
5 Evaluation	27
5.1 Evaluation protocol	27
5.2 Determining the baseline	29
5.3 Applying PSO to BBA	29
5.4 Towards distribution of the attack	31
5.5 Throwing the defense off the scent	39
5.6 Optimizing the attack	42
5.7 Comparing to state of the art	43
6 Conclusion	47
A Parameter table	51

Bibliography

59

List of Figures

2.1	Linear separability	5
2.2	Decision boundaries and decision regions	5
2.3	Activation functions	5
2.4	Example of a convolution layer	6
2.5	Difference targeted and untargeted attack	8
2.6	Adversarial training	10
2.7	PSO communication topologies	13
2.8	Particle swarm optimization	13
3.1	Difference between noise patterns	16
3.2	Influence of frequency on Perlin noise	17
3.3	Intuition of the Boundary Attack	18
3.4	Intuition of the HopSkipJumpAttack	19
3.5	Deep similarity encoder	20
5.1	Some examples of the MNIST dataset	28
5.2	Some examples of the CIFAR-10 dataset	28
5.3	Intuition of multiple starting points	32
5.4	Detections of the different attacks	32
5.5	Schematic overview of the query submission distribution	35
5.6	Overview of the modified round-robin based distribution scheme	35
5.7	Detections for more nodes	37
5.8	Detections for combinations of experiments	43
5.9	Visual comparison of the adversarial MNIST examples of different attacks	45
5.10	Visual comparison of the adversarial CIFAR examples of different attacks	46

List of Tables

5.1	Model architectures for the MNIST and CIFAR model	28
5.2	Model architectures for the MNIST and CIFAR similarity encoders . . .	28
5.3	Baseline results	29
5.4	PSO-BBA results	33
5.5	Distributed PSO-BBA results	36
5.6	Flushing-Embedded-Distance-Based distribution results	38
5.7	Results for different buffer sizes	38
5.8	The average inter query distances for different types of noise	40
5.9	Results for different buffer sizes	41
5.10	Comparison with state of the art	44
A.1	Parameter table	51
A.2	Parameter table of the defense	56

List of Abbreviations

AI Artificial Intelligence. 3

ANN Artificial Neural Network. 3, 4, 6

API Application Programming Interface. 7, 24, 25, 33, 57

BA Boundary Attack. 15, 16, 17, 18, 43

BBA Biased Boundary Attack. 16, 18, 21, 24, 29, 30, vii, 31, 32, 33, 36, 38, 41, 42, 43, 44, 45, 46, 52, 53

CNN Convolutional Neural Network. 3, 4

DB Distance-Based. 34, 36, 37, 52

DDoS Distributed Denial of Service. 21

DkNN Deep k-Nearest Neighbors. 9

DNN Deep Neural Network. 3

EA Evolutionary Algorithm. 9, 20

EDB Embedded-Distance-Based. 34, 36, 37, 38, 39, 40, 41, 52

FEDB Flushing-Embedded-Distance-Based. 37, 38

FGSM Fast Gradient Sign Method. 7

HSJA HopSkipJumpAttack. 17, 18, 36, 43, 44, 45, 46

MGRR Multi-Group with Random Redistribution. 11, 12

MRR Modified Round-Robin. 33, 34, 36, 52

PSO Particle Swarm Optimization. 9, 10, 11, 12, 13, 21, 24, 29, 30, vii, 31, 32, 33, 36, 38, 41, 42, 43, 44, 45, 46, 51

LIST OF ABBREVIATIONS

ReLU Rectified Linear Unit. 3, 5, 28

RNN Recurrent Neural Network. 3

RR Round-Robin. 33, 34, 35, 36, 42, 44, 52

Chapter 1

Introduction

Chapter 2

Background

2.1 Neural networks

Ever since the invention of computer systems, it has always been a goal of scientists and engineers to create **Artificial Intelligence (AI)**. Current state of the art approaches are mimicking the human brain, more specifically the neurons inside the brain. Already in the fifties, Rosenblatt introduced his perceptron [1]. The perceptron is a single neuron able to learn linearly separable patterns. It does so by finding a hyperplane that separates the two classes. This hyperplane is called the decision surface or decision boundary and the perceptron itself is called a classifier. Geometric regions separated by a decision boundary are called decision regions. The concept of linear separability is explained in Figure 2.1 in two dimensions. In Figure 2.2, the decision boundaries and decision regions are explained visually.

Unfortunately not all patterns are linearly separable. To overcome this problem, the neurons can be layered, creating an **Artificial Neural Network (ANN)** in the process. Layering neurons sequentially is essentially a linear combination of neurons. This in itself does not create non-linear decision surfaces. Non-linear activation functions are added for the **ANN** to be able to learn more complex decision boundaries. Some commonly used activation functions are **Rectified Linear Unit (ReLU)** [2], Heaviside step function [3] and softmax [4] (or sigmoid when used on scalars). In Figure 2.3 the plots of some activation functions can be found.

The neurons can be combined in different ways to create different **ANN** architectures. Each architecture has its own strengths and weaknesses. **Convolutional Neural Networks (CNNs)** excel in classifying visual data [5, 6], whilst **Recurrent Neural Networks (RNNs)** are widely used when there exist dependencies inside the data, such as in speech recognition [7, 8] or time series prediction [9]. More recent research focuses on **Deep Neural Networks (DNNs)**, due to the ever increasing computational power available. **DNN** approaches are able to compare to and even surpass human performance on very specific tasks [10, 11].

2. BACKGROUND

Most **ANNs** are not built from individual neurons, but rather from layers of neurons. An **ANN** architecture describes the different layers of which the network consists. The most basic type of layer is the fully connected or dense layer. This layer, as the name suggests, consists of neurons that are connected to all neurons of the previous layer. Every neuron in this layer outputs a value based on the perceptron rule:

$$y = b + W \cdot X$$

Here y is the output of the neuron, b is the bias, W is the weight matrix and X is the vector consisting of the outputs of the previous layer. The values of the bias and weights are learned based on the training data. $W \cdot X$ is the dot product $\sum_{i=1}^n w_i x_i$, where n is the number of neurons in the previous layer [12].

Activation layers apply activation functions to their inputs. They can be standalone layers, but more frequently they are integrated in other layers.

Convolution layers [13] are the main building blocks of **CNNs**. Convolution layers map inputs to so called feature maps. A filter W slides over the input X from left to right top to bottom. At every position, the dot product $W \cdot X$ of the frame and the underlying values of the input are computed. This value is placed at the corresponding position of the feature map. Higher values of the dot product (for a fixed value of $\|W\|_2$) mean that W and the underlying values of X are more similar. The feature map therefore indicates where in the input X a certain pattern W can be found. This process is translation-invariant, meaning that a certain pattern can be detected at any location of X . Scale and rotation invariance can be achieved by repeating the process for scaled and rotated filters. This is the main reason why **CNNs** are widely used for visual data. An example of a convolution layer is shown in Figure 2.4.

The use of convolution layers tends to cause an explosion in the dimensionality of the data. One filter will produce a feature map that is smaller than the original input. But a convolution layer typically consists of many filters, causing the explosion in the number of connections and weights. Pooling or subsampling layers [13] alleviate this problem by aggregating nearby points. The aggregation happens using a sliding window similar to a convolution layer. This approach of reducing the dimensionality has the added benefit that the output of the model is less sensitive to the exact location of a pattern. Some commonly used pooling layers are the MaxPooling and AveragePooling layers. They aggregate nearby points using the maximum and average value respectively. Convolution and pooling layers are often used in combination with each other and are repeated multiple times. This allows the **CNN** to detect patterns at different scales.




FIGURE 2.1: Linearly separable classes on the left and non-linearly separable classes on the right. Two classes are linearly separable if there exists a hyperplane for which all examples of one class are on the same side of this hyperplane, whilst all examples of the other class are on the other side of the hyperplane. In two dimensions, the hyperplane is a straight line.




FIGURE 2.2: Decision boundaries (blue lines) separate different decision regions (colored regions).




FIGURE 2.3: Plots of different activation functions. From left to right: ReLU, Heaviside step and sigmoid.




FIGURE 2.4: Example of a convolution layer. The filter slides over the input from left to right top to bottom. At every position, the dot product of the filter and input image is computed. This value is added to the corresponding position of the feature map. This process is repeated for m different filters leading to m feature maps.

2.2 Adversarial attacks

The expressiveness of ANNs is a double-edged sword. It is the cause for the near-human performance on some tasks, but also for counter-intuitive properties. As studied by Szegedy et al [14], one of these properties is the presence of discontinuous decision boundaries. This might cause seemingly identically images to be classified differently. They were the first to define adversarial examples as *imperceptibly small perturbations to a correctly classified input image, so that it is no longer classified correctly* [14]. This property of ANNs might not seem important at first glance, but it can be quite worrisome from a security point-of-view. Malicious users could craft images to bypass face recognition software [15] or attack the camera of a self-driving car to misclassify traffic signs [16]. Other fields where adversarial examples are of interest include malware detection [17], natural language processing [18] and industrial control systems [19]. Adversarial attacks are algorithms used to craft such adversarial examples.

All adversarial attacks are evaluated against a threat model. A threat model is a *structured representation of all the information that affects the security of an application* [20]. This information consists of the goals, knowledge and capabilities of the attacker, the accessibility of the model under attack and the costs of (un)successful attacks.

Most research on adversarial attacks is done using images. Researchers have the most freedom in this domain, since a slightly altered image is still an image with roughly the same contents. Slightly modifying an industrial control system however, might cause complete disruption of the system. Research in other domains is mostly conducted by altering existing image algorithms to the specific use case. For this reason this work focuses on adversarial attacks on images only.

2.2.1 Adversarial attacks terminology

Adversarial attacks are generally divided in two categories, white box attacks and black box attacks. In a white box attack, the attacker has complete knowledge of the classifier under attack. This knowledge consists of the architecture, parameters and thus their gradients and all output of the classifier. Examples of white box attacks are the **Fast Gradient Sign Method (FGSM)** [21] or the Carlini & Wagner attack [22].

In black box attacks, the only thing the attacker has access to is the output of the model. Depending on the literature, this output consists of class labels only (decision-based attack) or class labels and the corresponding confidence scores (score-based attacks). Black box attacks are more relevant in real-life scenarios, since most attacks are performed on a third-party **Application Programming Interface (API)**. These **APIs** generally do not reveal the underlying model.

Transfer attacks [23] try to overcome this hurdle by creating a surrogate model. This is an undefended model similar to the model under attack. This idea is based on the observation that adversarial perturbations often are transferable to other models [21]. Attacks can leverage information (such as gradients) from the surrogate model to breach the black box model. Due to the transferability of adversarial examples, it is also possible to perform so called zero-query attacks. Zero-query attacks are performed entirely on the surrogate model and the resulting adversarial example is forwarded to the black box model.

Both white box and black box attacks can be divided into targeted and untargeted attacks depending on their goal. In a targeted attack, the goal of the attacker is to create an adversarial example with a specific target class. In an untargeted attack the target class can be any class. Untargeted variants of attacks generally enjoy much more freedom and are therefore able to craft adversarial examples that are closer to the original. Figure 2.5 visually explains the difference between the two types.

What does it mean for images to be close to each other? This is easy to visualize in two dimensions as in Figure 2.5, but in higher dimensions, this is more difficult. Images reside in d -dimensional space, where d is the amount of pixels of the image. Two commonly used distances in higher dimensions are the L_2 -distance and the L_∞ -distance. They are defined as follows:

$$\begin{aligned} L_2(X, Y) &= \sqrt{\sum_{i=1}^d |x_i - y_i|^2} \\ L_\infty(X, Y) &= \lim_{p \rightarrow \infty} \left(\sum_{i=1}^d |x_i - y_i|^p \right)^{1/p} \\ &= \max(|x_1 - y_1|, |x_2 - y_2|, \dots, |x_d - y_d|) \end{aligned}$$




FIGURE 2.5: Different decision regions are shown in different colors. An adversarial example is being created starting from the source image (black star). On the left an untargeted attack is performed. The adversarial example is the image closest to the source image, that is classified differently (black circle). On the right a targeted attack is performed with the red decision region being the target class. The adversarial example is the image closest to the source image that is in the red decision region.

In both distances X and Y represent the images and $(x_1, x_2, \dots, x_i, \dots, x_d)$ and $(y_1, y_2, \dots, y_i, \dots, y_d)$ are the pixel values of X and Y respectively. The L_2 -distance is also known as the Euclidean distance, which is a generalization of the Pythagorean theorem in more than two dimensions. It takes the pairwise distances between all pixels into account. The L_∞ -distance is also called the Chebyshev distance. This distance only depends on the maximal pairwise distance between the two images. By minimizing the L_∞ -distance, the maximal pixelwise difference is minimized [24]. Sometimes the norm notation $\|\cdot\|$ is used instead of the distance notation. A distance is induced by a norm by $d(X, Y) = \|X - Y\|$. The other direction, a distance inducing a norm, is not as trivial. In order for a distance to induce a norm, the distance has to satisfy the following condition: $d(\alpha X, \alpha Y) = |\alpha|d(X, Y)$. This condition is satisfied in most commonly used distances, but not always. The discrete distance [25] is an example of such a distance and is defined as follows:

$$d(X, Y) = \begin{cases} 1, & \text{if } X \neq Y \\ 0, & \text{else} \end{cases}$$

2.2.2 Adversarial defenses

The existence of adversarial attacks naturally gave rise to adversarial defenses. These defenses can be categorized based on their objective. They can be either proactive or reactive. The goal of a proactive defense consists of making the models under attack more robust, while reactive defenses aim to identify attacks before they reach the model [26]. Different defensive countermeasures can be taken in each cat-

egory. The remainder of this section will discuss some commonly used techniques [27].

Gradient masking techniques hinder optimization-based attacks by having gradients "*that are not useful*" [28]. They are also sometimes referred to as obfuscated gradients [29]. Three types of obfuscated gradients can be identified. Shattered gradients introduce incorrect or non-existent gradients. Stochastic gradients are caused by random effects in the defense and exploding or vanishing gradients are primarily caused by chaining neural network evaluations. Besides intentionally introducing gradient masking in neural networks, they can also be introduced unintentionally due to the design of the network.

Defensive distillation [30] is a technique that can be classified as gradient masking. The goal of defensive distillation is to smooth the gradients of the model, making it more resilient to small input perturbations. The distillation procedure is done in two steps. First the model is trained on the original data and labels. This step produces a probabilistic output for each input due to the softmax activation function. Then the network is retrained using the original data and the probabilistic outputs as labels. The probabilistic labels contain additional knowledge that can be exploited to increase generalizability of the model.

Adversarial training [21] techniques inject adversarial examples in the training dataset and retrain the model in order to create a more resilient model. This is essentially a brute force method to correctly classify some adversarial examples. However this new model is still susceptible to new adversarial attacks, since the decision boundary has only been moved slightly. This can be seen in Figure 2.6.

Preprocessing techniques work on the inputs of a model. Different preprocessing techniques, such as denoising [31], dimensionality reduction [32] and image transformations [33] can be used to defend the model under attack. The goal of all techniques boils down to giving the attacker less control over the exact input that is being fed to the model.

Some defenses rely on **proximity measurements** between the input and the model. An example of this countermeasure is **Deep k-Nearest Neighbors (DkNN)** by Papernot and McDaniel [34]. **DkNN** computes support for a decision from a network based on a nearest neighbors search in the training data. Another example is region-based classification [35], where a prediction is made for a given input based on the proximity of training examples.

2.3 Particle swarm optimization

Particle Swarm Optimization (PSO) [36] is an optimization framework part of the **Evolutionary Algorithms (EAs)** family. In **EAs**, populations of candidate solutions




FIGURE 2.6: Decision boundaries before (left) and after (right) adversarial training. Before adversarial training, an adversarial example can be created (black circle). This example is added to the training set and the network is retrained. A new decision surface is created in the process. This surface classifies the previous adversarial example correctly (empty circle), but there is a new opportunity for an adversarial attack (black circle).

evolve based on mechanisms inspired by the field of biology, such as ant colonies [37], mutation and recombination [38]. The mechanism that inspired **PSO** is the behaviour of flocks of birds. The framework has been applied to numerous problems such as routing problems [39, 40], diagnosing diseases from imaging [41] and calculating heat transfer coefficients [42].

In **PSO**, different particles x move through the search space based on a set of rules. Their new position x_t is determined by their previous position x_{t-1} and a velocity v_t . The velocity depends on the distance to best position of the swarm g and the best known position of the particle p . The distances can be weighted by acceleration coefficients c to put more emphasis on exploration or exploitation. These values are multiplied by random parameters r uniformly distributed in $[0, 1]$. The velocity is also dependent on the previous velocity with a corresponding weight w . The best position is determined using a fitness function. This function states how ‘fit’ or good a certain position is with respect to the goal of the optimization problem. Equations 2.1 and 2.2 correspond to the update rules. Each particle i has its own position $x_{t,i}$ and velocity $v_{t,i}$ at time step t , but the i index has been omitted for readability. In Figure 2.8, the steps are graphically represented for a single particle.

$$v_t = \underbrace{wv_{t-1}}_{\text{Inertia}} + \underbrace{c_p r_p (p_{t-1} - x_{t-1})}_{\text{Individual best}} + \underbrace{c_g r_g (g_{t-1} - x_{t-1})}_{\text{Swarm best}} \quad (2.1)$$

$$x_t = x_{t-1} + v_t \quad (2.2)$$

Ever since the first mention of **PSO** in 1995, efforts have been made in order to improve the framework. The rest of this section will discuss some improvements.

The first version of **PSO** used a fixed value for the inertia weight w . Later it has empirically been shown that a linearly decaying weight improves performance [43]. This allows the algorithm to focus on exploration in the early iterations, while

shifting its focus to exploitation later on. The rate of decay depends on the high start value w_{start} , the lower end value w_{end} and the maximum number of iterations t_{max} . The weight in iteration t is calculated as follows:

$$w_t = w_{end} + (w_{start} - w_{end}) \left(1 - \frac{t}{t_{max}}\right) \quad (2.3)$$

In a large search space, particles can be far apart, which can cause the velocities to explode. By limiting the value of velocities to v_{max} , the swarm can be stabilized and the probability of finding an optimum is increased [36].

Clerc and Kennedy [44] studied **PSO** from a dynamic systems point of view. They provided a theoretically backed solution to the problem of the exploding velocities. Instead of limiting the value of the velocity, something which has only been shown empirically that it works, they proposed the constriction factor χ . The constriction factor is able to prevent the velocity explosion, whilst avoiding premature convergence to local optima. The velocity update of equation 2.1 is altered as follows:

$$v_t = \chi(v_{t-1} + c_p r_p(p_{t-1} - x_{t-1}) + c_g r_g(g_{t-1} - x_{t-1}))$$

The inertia weight w is dropped and the entire sum is multiplied with χ . The value of χ is calculated as follows:

$$\chi = \begin{cases} \sqrt{\frac{2\kappa}{C-2+\sqrt{C^2-4C}}}, & C > 4 \\ \sqrt{\kappa}, & \text{else} \end{cases}$$

Here C is the sum of the acceleration coefficients c_p and c_g and κ is a user defined value to determine the rate of convergence.

Another improvement on vanilla¹ **PSO** is **Multi-Group with Random Redistribution (MGRR) PSO** [46]. The swarm is split into two groups with opposite acceleration coefficients. The first group, with a larger c_p , is more attracted to the local optima, while the second group, with larger c_g , moves to the global optimum. The combination of the two groups, accounts for more variation during the search process. This in turn aids in the convergence to the global optimum.

Even with the multi-group improvement, the particles can get trapped in local optima, as is the case in vanilla **PSO**. Whenever the swarm is stuck in a local optimum, half of the particles are randomly redistributed. The particles that will be redistributed will change approximately half of their values. The redistribution gives the algorithm a chance to escape the local optimum.

Experimental results show that **MGRR-PSO** has better performance than the vanilla variant [46]. The convergence is also faster and less dependent on the starting

¹Vanilla is used to refer to software or algorithms not altered from their original form [45].

positions of the particles. Even the variant without the random redistribution outperforms vanilla **PSO**.

PSO is often combined with other algorithms in so called hybridization techniques. The goal of this combination is to create an algorithm that contains the beneficial properties of both its constituents. Some commonly used algorithms combined with **PSO** are genetic algorithms [38], differential evolution [47] and ant colony optimization [37]. The hybridization is most commonly done in three different ways. In the first way, one algorithm is used before the other. The first algorithm performs optimization on the population used in the second algorithm. The second way divides the entire population into two groups and every group is optimized with a constituent algorithm. Some inter group communication is required for this approach to work. The idea behind this technique is similar to the **MGRR** approach discussed earlier. The third way incorporates specific parts of an algorithm as a local search operator in the other. Thangaraj et al. [48] studied 64 hybridization algorithms containing **PSO** as one of the constituents. They found that the combination of differential evolution and **PSO** edges out the other hybrid approaches in terms of average error. However, all techniques outperform vanilla **PSO** using the error metric.

Another approach to prevent the particles of the swarm to converge to local optima is to restrict the communication inside the swarm. In vanilla **PSO** all particles have knowledge about the best position of the entire swarm. All particles will take a step towards this position in the next iteration of the algorithm. This causes the particles to converge to each other and limits the exploration of the search space. In order to restrict this knowledge, different communication topologies can be used. Some topologies are depicted in Figure 2.7. Here the nodes represent the different particles and the edges show the possibility of communication between these particles. The fully connected topology is the standard topology used in **PSO**. A ring topology limits communication of the best position to two neighbors per particle. The pyramid topology is similar to the ring topology, but there exists one particle that can communicate to all other particles. Neighbors in the topologies can be determined based on pre-defined indices or on the position of the particles in search space. The latter requires more computations. It has been shown that the best topology is highly problem dependent and that the size of the swarm plays a significant role in the choice of the best topology [49]. Combinations between the fully connected and ring topology are also possible and show promising results [50].




FIGURE 2.7: Different communication topologies for the **PSO** framework. The particles are represented by the nodes and the edges represent the communication channels. Inspired by [49].




FIGURE 2.8: Particles move based on a step towards their best known position, the swarm's best known position and a step in the direction of the movement of the previous iteration. The different steps are combined to determine the new position of the particle.

Chapter 3

Related work

3.1 Boundary attack

Boundary Attack (BA) [51] is a decision-based adversarial attack. The basic intuition of **BA** differs from traditional adversarial attacks. Unlike these traditional adversarial attacks, where the original image is moved through search space in order to become adversarial, **BA** starts from an input that is already adversarial. This input is then moved closer to the original image, while staying adversarial.

The attack has to be initialized with an already adversarial input. Two different approaches can be taken depending on the attack setting. In the untargeted case, the input can be sampled from a maximum entropy distribution given the valid domain of this input. Samples that are not adversarial are rejected. An example of such a starting position can be seen in Figure 3.1a. In the case of a targeted attack, the input is a sample from the dataset that is classified as the target class by the model under attack.

BA iteratively updates the adversarial image by performing a step orthogonal to the original image and a step towards this image. In iteration k , a perturbation η_k is sampled from a uniform distribution. This perturbation is rescaled and added to the adversarial image. From this new position in search space, the step towards the original image is taken. This way the path of the attack follows the decision boundary, hence the name of the attack. The intuition of the **BA** is shown in Figure 3.3. The attack can only follow the boundary if the adversarial image is already near the boundary. The starting image is projected onto the boundary using binary search to ensure that the adversarial image is in the vicinity of the boundary.

The step sizes are adjusted according to local geometry of the boundary. The orthogonal step size δ is adjusted so that approximately half of the orthogonal perturbations is still adversarial. This approach is based on trust region methods [52]. The step size towards the original image ϵ is adjusted using the same principle, but here a user specified threshold is used. The decision boundary tends to become flatter, the closer to the original image the attack gets [53]. Therefore the algorithm converges




FIGURE 3.1: Difference between noise patterns.

when ϵ converges to zero.

Biased Boundary Attack (BBA) [54] (previously known as *Boundary Attack++*) is an improvement on the original **BA** in three different ways. All three improvements will be discussed in order of the strength of their effect. The first improvement is a biased sampling technique. The key idea behind this strategy is that most previous attacks yield adversarial examples with high frequencies in the image. By sampling the perturbations in the first step of the **BA** from a low frequency distribution, the frequency of the created adversarial example will be lowered as well. **BBA** does this by sampling from a Perlin noise [55] distribution instead of a uniform distribution. Lower frequency images yield more natural results and can more easily bypass simple preprocessing defense schemes. The difference between the two noise patterns can be seen in Figure 3.1. The noise patterns can be influenced by a frequency value. This value can be tuned depending on the size of the images at hand. Higher frequency values yield less smooth noise patterns. Figure 3.2 visually shows the influence of the frequency values.

The second improvement is to use a regional mask. The original **BA** applies a perturbation to the images as a whole. Every pixel will be perturbed with the same magnitude. This magnitude can be altered on a per-pixel basis when using a mask. Pixels that are further away from the target image will receive a larger perturbation than pixels that are already close to the corresponding pixel in the target. The mask m is constructed according to equation 3.1 based on the original image x_{orig} and the adversarial image x_{adv} . It is then pixel-wise applied to the sampled perturbation in equation 3.2. The masked perturbation is normalized afterwards.

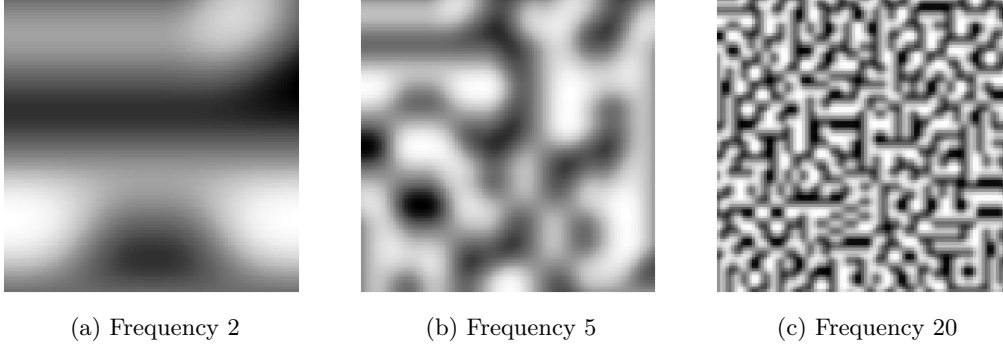


FIGURE 3.2: Influence of frequency on Perlin noise patterns.

$$m = |x_{adv} - x_{orig}| \quad (3.1)$$

$$\eta_k = m \odot \eta_k; \eta_k = \frac{\eta_k}{\|\eta_k\|} \quad (3.2)$$

This technique improves efficiency since the search space is significantly reduced. It is also possible to engineer masks for specific examples in order to incorporate other knowledge in the attack.

The final improvement is based on the idea of transfer attacks. A surrogate model is trained and will be used to calculate adversarial gradients. These gradients will then be used to bias the sampling direction for the orthogonal step. If the surrogate model does not closely resemble the defender, then the gradients will only hamper the speed of convergence of the attack instead of causing the attack to fail.

3.2 HopSkipJumpAttack

HopSkipJumpAttack (HSJA) [56], like **BA**, is a decision-based adversarial attack that starts from an adversarial input. The initial input is obtained in an identical manner as in **BA**. **HSJA** is an iterative algorithm that consists of three steps.

The first step is a projection onto the decision boundary of the model under attack. This projection is carried out using a binary search. The second step is to estimate the direction of the gradient at the boundary. Different directions are sampled from a uniform distribution over a d -dimensional sphere, where d is the input dimension. This random direction is added to the boundary point, generating a new query for the model. The results of these queries are combined to a gradient estimation $\bar{\nabla}S$ using the Monte Carlo estimate of equation 3.3. In this equation u_b are the random directions and x_t is the boundary position. B is the number of random directions that needs to be sampled. This number increases based on the current iteration of




FIGURE 3.3: Intuition behind the Boundary Attack. On the left the path of the attack is shown. The first step is a projection onto the boundary, afterwards it follows the decision boundary of the class of the original image. Each arrow represents one iteration of the attack. On the right, the two different steps of each iteration can be seen. In the first step, a random direction is sampled and projected onto a sphere around the original image. The second step is to take a step towards the original image from this new position. Image inspired by [51].

the attack to reduce the variance of the estimate. The function ϕ_{x^*} returns 1 if the new position is adversarial and -1 if it is not adversarial. δ is a positive parameter determining the size of the d -dimensional sphere.

$$\widetilde{\nabla S}(x_t, \delta) := \frac{1}{B} \sum_{b=1}^B \phi_{x^*}(x_t + \delta u_b) u_b \quad (3.3)$$

Once the gradient has been estimated, the third and final operation is to take a step along this gradient. The step size is determined using a geometric progression scheme. These steps are iteratively repeated until the pre-set stopping criterion is met. Figure 3.4 represents the intuition behind **HSJA** in a graphical manner.

HSJA eclipses **BA** and **BBA** both on median distance against queries and attack success rates using a limited amount of queries. The untargeted version of **HSJA** is able to compete with white box attacks on the ImageNet dataset [57]. It also performs similar or superior to white box attacks such as the C&W attack [22] when evaluated against defensive mechanisms such as defensive distillation [30], region-based classification [35] and adversarial training [21].




FIGURE 3.4: Intuition behind the HopSkipJumpAttack. Each iteration consists of three steps. The first step is a projection onto the boundary. The second step is the estimation of the gradient at this point. This is done by sampling directions from a uniform distribution and querying the model under attack from this new position (grey arrows). The results are combined via the Monte Carlo estimate. The third and final step is to take a step along the estimated gradient. Image inspired by [56].

3.3 Stateful defense

The defensive schemes discussed in section 2.2.2 all operate on the query level. They try to detect and flag possible attacks based on a single query without taking other context into account. The stateful detection mechanism by Chen, Carlini and Wagner [58] is different in this aspect. As the name suggests, it holds state of previously submitted queries. It is similar to the defenses that use proximity measurements, but the measurement is between queries instead of between the query and training data.

All queries submitted to the model equipped with a stateful detection mechanism are stored in a history buffer. Each user of the model has a distinct history buffer, where its queries are stored. These buffers can be bounded by time or number of queries depending on the resources available and the use case of the model. Each time a query is submitted to the model, the average distance to its k nearest neighbors is calculated and if this distance is lower than a certain threshold, then the user gets flagged by the mechanism. Appropriate actions such as banning the account can be taken.

The distance metric is not calculated in input space. Each query is encoded by a deep similarity encoder [59] to an encoded space, typically of a lower dimension. In this encoded space, images which represent perceptually similar objects are clustered together. A visual of the idea behind the similarity encoder is shown in Figure 3.5. The advantage of the encoded space is twofold. Firstly, the dimension of the encoded




FIGURE 3.5: Visual representation of a deep similarity encoder. Input images (on the left) are mapped to an embedded space (of lower dimension). Images representing the same concept (squares, circles and triangles) are mapped close to each other in clusters. A nearest neighbor search around the embedding of an image will therefore yield similar images. The example images are taken from the MNIST dataset [13].

Image inspired by [59].

space is smaller than the dimension of the input space. Therefore less space is needed to store the history buffers. Secondly, simpler distance metrics such as L_2 -distance in input space can easily be evaded by an attacker. For example the L_2 -distance can be significantly increased by simply rotating or shifting the input image.

The parameter k , the number of neighbors to consider is picked as follows. As the training data of the model consists of only benign queries, no attacks should be flagged when feeding the stateful detection mechanism with this data. To allow for some more leniency, a false positive rate of 0.1% is still acceptable. For each value of k , a different threshold will be required to maintain the selected false positive rate. Larger values have the benefit of larger thresholds causing the defense to be more resilient, since attackers' images need to be more diverse. But k is also the number of queries needed before an attack can be flagged. Therefore too large values for k are disadvantageous. Smaller values also reduce computational cost. Chen, Carlini and Wagner set the value of k to 50 for the CIFAR-10 dataset [60], since the thresholds increased sharply up to this value. Other datasets might require different values for k .

3.4 PSO and distributed attacks

There have been several attempts to craft adversarial examples using an EA. Previous attempts tried to reduce the number of queries needed to create a successful adversarial example by utilizing EAs [61, 62, 63, 64, 65, 66].

GenAttack [61] and the similar efficient attack by Dong et al. [62] use genetic

algorithms in order to minimize the number of queries to the model. Both algorithms reduce the dimension of the search space to improve the efficiency of the attack. Once a promising perturbation is found in this lower dimensional space, it is upscaled using a bilinear transformation. By reducing the search space, the number of individuals in the genetic algorithm can be lowered, which in turn lowers the total amount of queries. GenAttack also uses annealing schemes to adaptively scale the parameters of the algorithm. This allows it to escape local optima and improve the adversarial example further.

AdversarialPSO [63] and the similar attack from [64] use **PSO** as optimization routine on images and audio fragments respectively. Each particle represents a possible adversarial example. Both attacks use the standard rules of **PSO** as specified by [36] improved with a linearly decaying inertia weight [67]. The former attack also uses a constriction factor to avoid premature convergence [44]. While the latter solves this problem by generating new particles using a genetic algorithm when premature convergence is detected. Both algorithms rely on confidence scores to assign fitness values to certain positions in the search space. **PSO-BBA** [65], is similar to AdversarialPSO, but only relies on distances to determine fitness values. This attack can therefore also be used in decision-based settings instead of solely in score-based settings.

The idea behind the multi-group **PSO** attack [66] is to use multiple **PSO** swarms to escape local optima. The intuition behind it is inspired by the **Distributed Denial of Service (DDoS)** attack [68]. The swarm is split into multiple smaller groups and each group is placed on a single node. The groups perform the standard **PSO** algorithm. The best position over all groups is communicated using a dedicated server. Each group submits its queries from its own node, tricking the defensive mechanism into thinking that multiple users are submitting queries. The authors state that this will ultimately result in less detections, but they have not evaluated this against a defensive scheme.

Chapter 4

Approach

The research concerning adversarial attacks and defenses is predominantly driven by a game of cat and mouse. Whenever a new attack is proposed, a defensive mechanism countering the novel attack is developed and vice versa. This work aims to create a new family of algorithms that can be used to perform both targeted and untargeted attacks. The goal of these algorithms is to craft adversarial examples comparable to state of the art approaches while remaining undetected by the stateful detection mechanism [58] described in section 3.3.

4.1 Distribution

The stateful detection mechanism [58] makes the assumption that queries can be traced back to their adversary and that there is no cooperation between different adversaries. This assumption can be problematic as N collaborating adversaries can theoretically reduce the number of submitted queries per adversary by a factor $1/N$. Even a single adversary could set up multiple accounts and submit queries on each account until it is banned. Due to the reduced number of submitted queries, less attacks will be detected since each buffer of the defense mechanism only holds a fraction of all queries.

This work will aim to evade the detection mechanism by distributing the query submissions over multiple nodes. Each node will represent a different user of the model under attack. The users could theoretically be all different persons or they could be different accounts of the same person.

As described in section 3.4, several attempts have been made to distribute adversarial attacks [65, 66]. However, none of these attacks have been evaluated against the stateful detection mechanism, since the goal of the distribution was to make the attack more efficient in terms of the distance between the original image and the resulting adversarial example. This work will distribute the query submission over multiple nodes in order to avoid detection.

4.2 Optimization

As previously mentioned, reducing the number of submitted queries per adversary by a factor $1/N$, where N is the number of collaborators, is straightforward. Adversaries can gain knowledge about the search space by cooperating with other adversaries. They can leverage this knowledge in order to reduce the number of submitted queries even more. This idea has been utilized by multiple algorithms that were mentioned in section 3.4. These algorithms used some form of **PSO** to optimize the final adversarial example. However all but the **PSO-BBA** algorithm by Xiang et al [65] rely on the confidence score of the model for the fitness value calculation. All attacks discussed use **PSO** as an attack in itself. This work will combine the benefits of state of the art black box attacks and **PSO**.

4.3 Approach

The remaining sections of this work will propose a new family of adversarial attacks. First a threat model is defined in section 4.4. All remaining experiments will be performed with this threat model in mind. Afterwards the novel adversarial algorithm is proposed in section 5.3. This algorithm is iteratively improved based on the results of the experiments. Finally, the final algorithm will be compared with state of the art decision-based attacks.

During the process of creating, improving and optimizing the attack, this work tries to answer the following research questions.

- What are the (dis)advantages of using **PSO** in relation to vanilla adversarial attacks?
- How can **PSO** be combined with state of the art adversarial attacks?
- What are the (dis)advantages of distributing an adversarial attack?

This work concludes with an answer on these questions in section 6.

4.4 Threat model

This section will describe the threat model that will be used for the remainder of this work. The first and most impactful assumption is that the model under attack will only output labels for the input. There will be no confidence scores or model parameters available. The proposed attack will have to be a decision-based attack. Most real world **APIs** will only expose the final decision to the user, causing decision-based attacks to be the only viable option. Decision-based attacks are the most restricted type of attack. Therefore they can also be applied to score-based and white-box models.

The proposed attack will be a targeted attack, as it is the most relevant type from a security point of view. A targeted attack implementation can also easily be ported to an untargeted one. This is done by running a targeted attack for every possible class and selecting the one closest to the original image.

The model will be defended by a stateful detection mechanism [58] as described in section 3.3. It will have a query bounded buffer for each account. Once an query has been flagged as potential attack, the account that submitted the query will be banned. The user will have to set up a fresh account in order to submit queries again.

There are two main fees associated with the model. The first fee is related to setting up an account. The assumption is made that setting up an account requires a valid credit card or phone number. Whenever an account is banned, the credit card or phone number is invalidated in the system, requiring the attacker to sign up for a new credit card or register for a new phone number. This cost is identical to the assumption made by Chen et al [58]. The second fee is related to the amount of queries submitted to the model. The more queries an attacker submits, the higher the total cost will be. Chen et al [58] did not incorporate a cost per query, but most vision APIs, such as Google Cloud Vision [69] and Amazon Rekognition [70] use this type of fee.

These fees ensure that attackers need to be both efficient and evasive in order to be successful. The evasiveness of the attack is correlated to the number of detections by the stateful defense mechanism. Each detection essentially means that the attacker has to set up a new account and pay the associated cost. To compare the efficiency of different attacks, the following strategy will be used. Every run of an attack has a budget of queries. The attack that has created an adversarial example closest to the original image inside this budget of queries, is the most efficient. The evasiveness of an attack is the primary concern, since the cost of setting up a new account is significantly higher than the cost per query.

Chapter 5

Evaluation

This chapter will discuss the experiments performed, as well as the ideas and intuitions behind them. The adversarial algorithm will be refined throughout the different subsections and the results of the refinement will be discussed at the end of each subsection.

5.1 Evaluation protocol

This subsection describes the evaluation protocol that will be followed for all experiments performed. Experiments will be performed with the MNIST [13] and CIFAR [60] datasets. Some examples of these datasets are visualised in Figures 5.1 and 5.2 respectively. A black box model is trained using the training data of the respective dataset. The architectures of the models are identical to the ones used in [22, 30]. A summary of the two architectures can be found in Table 5.1. The training parameters are also identical to the ones used in [22, 30]. These models remain unchanged for all experiments.

The trained models are then used to classify all instances in the test set of their respective dataset. The incorrectly classified examples are filtered out of this set as they are essentially already adversarial. A list of experiments is generated based on the remaining examples. An experiment consists of an original image, a target label and starting position(s). All future refinements will therefore perform the same set of experiments in order to make the comparison more fair. All random effects present in the algorithm are seeded for the same purpose.

The stateful defense mechanism [58] will use a query bounded detector buffer of 1000 queries per user and the value of k will be set to 50, as suggested by the authors of the paper. A threshold is determined in order to achieve a 0.1% false positive detection rate on the training data. The detector buffer will be cleared after each detection to simulate the creation of a new account. The mechanism will have a similarity encoder with an output dimension of 256. The detailed architecture of the



FIGURE 5.1: Some examples of the MNIST [13] dataset.




FIGURE 5.2: Some examples of the CIFAR-10 [60] dataset.

TABLE 5.1: Model architectures for the MNIST and CIFAR models. The architectures are identical to [22, 30].

Layer type	MNIST Model	CIFAR Model
Convolution + ReLU	$3 \times 3 \times 32$	$3 \times 3 \times 64$
Convolution + ReLU	$3 \times 3 \times 32$	$3 \times 3 \times 64$
Max Pooling	2×2	2×2
Convolution + ReLU	$3 \times 3 \times 64$	$3 \times 3 \times 128$
Convolution + ReLU	$3 \times 3 \times 64$	$3 \times 3 \times 128$
Max Pooling	2×2	2×2
Fully Connected + ReLU	200	256
Fully Connected + ReLU	200	256
Softmax	10	10

TABLE 5.2: Model architectures for the MNIST and CIFAR similarity encoders. Both datasets use the same architecture.

Layer type	Size
Convolution + ReLU	$3 \times 3 \times 32$
Convolution + ReLU	$3 \times 3 \times 32$
Max Pooling	2×2
Convolution + ReLU	$3 \times 3 \times 64$
Convolution + ReLU	$3 \times 3 \times 64$
Max Pooling	2×2
Fully Connected + ReLU	512
Fully Connected + linear	256

TABLE 5.3: Results for the **BBA** baseline approach on MNIST and CIFAR-10 dataset.

Attack	MNIST		CIFAR	
	Distance	Detections	Distance	Detections
Baseline BBA	2.807	413	1.306	474

encoder is shown in Table 5.2 and is identical for both datasets.

The algorithm receives a budget of 25 000 queries to create an adversarial example. The final adversarial example will be evaluated on the L_2 -distance to the original image and the total number of detections by the stateful defense mechanism.

5.2 Determining the baseline

The goal of this subsection is to determine a baseline to which all future algorithms will be compared. The baseline will be a vanilla **BBA** using the same hyperparameters as suggested in [54]. The orthogonal step will be set to 0.05 and the source step to 0.002. Both the Perlin noise improvement and the regional masking improvement will be used. No gradients of surrogate models will be calculated as the authors of the paper already mentioned that the improvement was marginal.

Table 5.3 reports the average L_2 -distance and the average number of detections for the different datasets.

5.3 Applying PSO to BBA

As mentioned in section 4.2, previous attempts to craft adversarial examples using **PSO** used **PSO** as a standalone algorithm to guide adversarial examples through the search space closer to the original image. This work aims to combine the benefits of **PSO** and an existing decision-based attack such as **BBA**.

Vanilla **BBA** has the disadvantage that, depending on the starting position, it might get trapped in a local optimum. By using **PSO** in combination with **BBA**, multiple starting positions can be explored and the probability of getting trapped is lowered. The intuition behind this idea is shown in Figure 5.3. The more starting positions there are, the higher the probability of finding the global optimum. There is however a clear trade-off in terms of efficiency. By having n different starting positions, the query budget is essentially reduced by a factor n for each starting position.

The efficiency reduction does not have to pose a problem due to the implicit communication in the swarm. Particles can move to promising regions in the search space based on the information of their peers. The promising regions are therefore more

queried.

The proposed **PSO-BBA** algorithm works as follows. Particles will perform a more aggressive version of **BBA**. The initial source step ϵ is set significantly higher than in the vanilla version. The increased source step might cause the particle to end up in a non-adversarial decision region. Once this happens, the standard **PSO** equations (2.1 and 2.2) are used to guide the particle back to the adversarial region. Later refinements of this algorithms will deal with attack iterations. An attack iteration is defined as a specified budget of the aggressive **BBA** (50 queries in all experiments) or a single query submission using the **PSO** rules.

The inertia weight of the **PSO** equations is determined by the linearly decaying scheme of equation 2.3, with the weight decaying from 1 to 0. The acceleration coefficients c are set based on an idea of multi-group **PSO**. Two groups with opposite acceleration coefficients are created. This approach helps escape local optima [46]. The equations for both c_p and c_g are:

$$c_p = \begin{cases} \max(A1, A2), & \text{if } i \bmod 2 = 0 \\ \min(A1, A2), & \text{else} \end{cases} \quad (5.1)$$

$$c_g = \begin{cases} \min(A1, A2), & \text{if } i \bmod 2 = 0 \\ \max(A1, A2), & \text{else} \end{cases} \quad (5.2)$$

Here i is the index of the particle in its swarm. The values of $A1$ and $A2$ are 1 and 2 respectively. These values are suggested by the authors of [66].

The source step ϵ will be changed after every iteration of the attack. Two separate multipliers are used to respectively increase and decrease the value of this parameter. The value of ϵ is slightly increased if the new position is still adversarial. Likewise it is decreased if the position is no longer adversarial.

By using **PSO** in combination with **BBA**, the advantage of multiple starting points, as explained in Figure 5.3, can be exploited, without having a less efficient attack as a whole. Whenever particles end up in non-adversarial regions, they will move closer to the best known position in the swarm due to the **PSO** equations. At the end of the attack, most particles will be in the same area of the search space, allowing for more exploitation in this specific area.

The **PSO** framework requires a fitness function to quantify the fitness of a position for the problem at hand. The authors of AdversarialPSO [63] suggested the following fitness function f :

$$f(x) = |p_x - p_{x'}| - \frac{c}{n} \|x - x'\|_2 \quad (5.3)$$

Where x is the position of the particle, x' is the original image, p_x and $p_{x'}$ are the confidence scores of the model in predicting the label of x and x' respectively and c is a constant to weight the penalty. However, as discussed in section 4.4, the confidence scores are not available for this specific attack. The fitness function in equation 5.3 also assumes that the position x will always be adversarial. This will not be the case in the proposed algorithm. The fitness function will therefore be altered to the following:

$$f(x) = \begin{cases} \|x - x'\|_2, & \text{if } x \text{ is adversarial} \\ +\infty, & \text{else} \end{cases} \quad (5.4)$$

The infinite value for the fitness function inside non-adversarial decision regions acts as a penalty, causing the particles to quickly diverge from these regions.

The same set of experiments as in section 5.2 has been performed. The experiments have been done using attacks with both five and ten particles. The initial step sizes have been set to 0.25 and 0.20 for MNIST and CIFAR respectively. The values for the increasing and decreasing multiplier have been set to 1.05 and 0.99 for both datasets. These values have been selected at random and were tested on a small subset of the experiments in order to confirm their effectiveness. Other values will be tested in section 5.6. The results of the experiments can be found in Table 5.4.

The five particle **PSO-BBA** algorithm outperforms the baseline in terms of distance to the original image on both datasets. This is not the case for the ten particle version. The high number of particles requires sufficient queries in the beginning of the attack in order to discover promising regions in the high dimensional search space of CIFAR. This exploration requires more queries than the query budget allows. The number of detections is lower for all variants of **PSO-BBA** compared to the baseline. The different starting points have the added advantage that initial queries are more spread out over the search space. These queries are therefore less similar and the detector will not flag as much attacks. This effect can be seen in Figure 5.4 for the MNIST experiments.

The number of detections drops as the number of particles increases. This was to be expected from the intuition behind **PSO-BBA**. However, the average L_2 -distance to the original image is higher for ten particles compared to five. To reduce the number of different parameters in future refinements, only swarms with five particles will be considered from here on.

5.4 Towards distribution of the attack

The stateful defense mechanism [58] makes the assumption that there is no collaboration between adversaries. Based on this assumption, each account or user will have its own detector buffer. This subsection aims to exploit this assumption by




FIGURE 5.3: The vanilla version of **BBA** might get stuck in a local optimum depending on the starting point (left plot). By starting from multiple positions, the probability that **BBA** gets stuck in a local optimum is reduced (right plot). The multiple starting points are particles in a **PSO** swarm. Image inspired by [51].

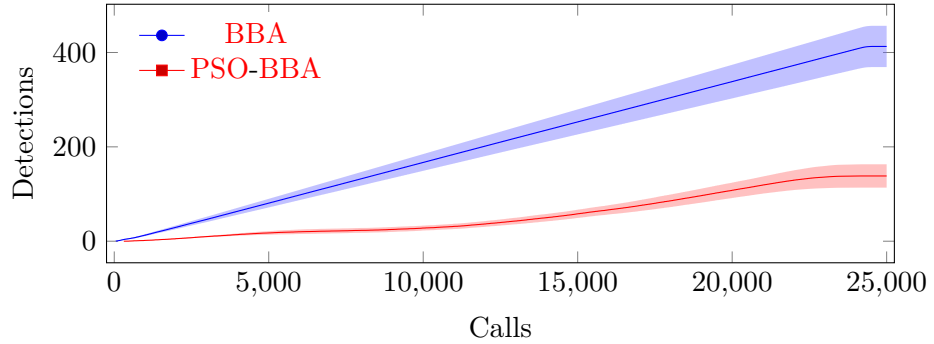


FIGURE 5.4: The cumulative number of detections for different attack algorithms for the MNIST experiments. The number of detections of the **BBA** algorithm steadily increases with the number of calls. The number of detections of the **PSO-BBA** algorithm increases more slowly in the beginning due to the dispersed nature of the attack. The increase is more sharp near the end of the attack when the particles converge. The plot shows the 95% confidence regions around the graphs.

TABLE 5.4: Results for the **PSO-BBA** approach on MNIST and CIFAR-10 dataset.

Attack	MNIST		CIFAR	
	Distance	Detections	Distance	Detections
Baseline BBA	2.807	413	1.306	474
PSO-BBA (5 particles)	2.712	138	1.239	257
PSO-BBA (10 particles)	3.157	44	2.290	184

distributing the query submission over multiple accounts.

By distributing the query submissions over multiple accounts, the number of total detection should be reduced compared to **BBA** due to two reasons. The first reason is obvious. If a budget of 25 000 queries is distributed over N accounts, then every account will submit $25\,000 / N$ queries. The less queries there are submitted, the less potential attacks will be detected. The second reason is due to **PSO**. As stated before, the different particles reside in different parts of the search space. This means that a detector buffer of a specific account will contain queries from all over the search space, causing the inter query distances to be larger. The latter reason was also present in the algorithm described in section 5.3.

It should be noted that there is only distribution at the query submission level. The algorithm itself will be executed on one machine without the need for distribution. Even the query submission can be done from one machine by constantly changing the **API** key or by logging in and out of a specific account whenever it is needed. This approach makes the assumption that the **API** under attack does not track the IP address of the machine that submitted the query. Multiple machines might therefore be more convenient in a real attack setting.

This work simulates the multiple machines setting by having separate node (or machine) objects on the same machine. Instead of having a detector buffer for each node at the location of the model, the responsibility is moved upstream to the nodes themselves. Each node will have a buffer to which the query will be added. All queries will be passed to the model afterwards. In Figure 5.5, the distribution is shown schematically.

The queries will be distributed over the nodes based on a distribution scheme. Three different distribution schemes will be used in the experiments. The first two schemes are inspired by the **Round-Robin (RR)** scheduling algorithm [71] and will be called the **RR** and **Modified Round-Robin (MRR)** distribution schemes. The third distribution scheme will be built on the assumption that the inter query distance in the detector buffers needs to be maximized.

The **RR** distribution scheme is fairly straightforward. The outgoing queries of the

algorithm will be evenly distributed over the nodes in circular order. This scheme does not have a notion of the underlying attack. This can cause successive queries forwarded to the same node to be relatively close together by chance.

The **MRR** distribution scheme maps a particle to a node for the entire duration of an attack iteration. The queries submitted by a single particle during one attack iteration are relatively close together, but this approach ensures that all particles will submit queries to all nodes. This might ultimately cause less detections. The **MRR** distribution scheme maps a number of particles on a number of nodes. If the amount of particles is equal to the number of nodes then the mapping is straightforward. If the number of nodes is less than the number of particles then dummy nodes are introduced in the mapping. The dummy nodes forward their received queries to a random existing node. The random node changes every rotation. If the number of particles is less than the number of nodes then dummy particles that do not submit queries, are introduced. After every attack iteration, the mapping rotates. The mapping process is shown in Figure 5.6.

The ultimate goal of the attack is to evade the stateful detection mechanism that defends the model under attack. Potential attacks are flagged based on the inter query distance in a detector buffer. The distance-based distribution scheme tries to exploit the defense. Instead of distributing queries over nodes in a **RR**-fashion, queries are submitted to nodes based on the distance to the previously submitted queries to this node. The node with the highest average distance to the previously submitted queries is selected. For this purpose, every node will hold an internal buffer of queries. The size of this buffer can be chosen by the attacker, as well as the distance metric. Standard L_2 -distance can be chosen, but this has the same pitfalls as discussed in section 3.3. It is also possible to train another similarity encoder and use the L_2 -distance in the embedded space in order to resemble the defense more closely. The distribution schemes will be called **Distance-Based (DB)** and **Embedded-Distance-Based (EDB)** distribution schemes respectively.

The experiments performed in this section will use a buffer size at each node of 20 for both the **DB** and **EDB** distribution schemes. The similarity encoder in the case of the **EDB** scheme will be trained on the test data of the respective dataset and has an output dimension of 128.

The results shown in Table 5.5 reveal that the different distribution schemes have similar performance in terms of the number of detections. The distance is identical for all distribution schemes since the underlying algorithm is the same. They only differ in the way that the queries are distributed over the nodes. The ten node distribution is always better than the five node version. An attacker has to make a trade-off between the number of nodes and the number of detections (and consequently accounts). More nodes require trigger less detection and therefore require less accounts. Depending on the cost of nodes and accounts, the optimal amounts might differ.




FIGURE 5.5: Schematic overview of the query submission distribution. Each arrow represents a query. The algorithm distributes the queries over the different nodes. Every node forwards its received queries to the model using the corresponding account. In a real setting, the model will have a detector buffer for every account.

In this work the buffers are located on the nodes themselves.




FIGURE 5.6: Overview of the modified round-robin based distribution scheme. The particles are represented by the small circles and the squares represent the different nodes. If the number of particles is equal to the number of nodes (left), then particles submit their queries to the corresponding node. If the number of nodes is less than the number of particles (middle), then the particles without nodes will submit their queries to a random node. This random node will change every attack iteration. If the number of particles is less than the number of nodes (right), then dummy particles are introduced. These dummy particles will not submit any queries, therefore some nodes will not receive queries. After every attack iteration the particle mapping will rotate.

As can be concluded from Table 5.5, there is no real gain in having a complicated distribution scheme. The **RR** distribution scheme will consequently be used for the remaining experiments since the distance based schemes have more computational overhead. These experiments will also be performed using ten nodes, unless mentioned otherwise.

TABLE 5.5: Results for the distributed **PSO-BBA** approaches on MNIST and CIFAR-10 dataset. There is no real advantage of one distribution scheme over the others. Increasing the number of nodes however, does have a clear advantage. The distribution over ten nodes is flagged less than the five node distribution for all schemes. The **DB** scheme performs best for the MNIST dataset, while the **RR** scheme is the best performer for the CIFAR dataset. Their respective values are shown in bold in the table.

Attack		MNIST		CIFAR	
		Distance	Detections	Distance	Detections
Baseline	BBA	2.807	413	1.306	474
	PSO-BBA	2.712	138	1.239	257
RR-PSO-BBA	(5 nodes)	2.712	124	1.239	224
MRR-PSO-BBA	(5 nodes)	2.712	110	1.239	227
DB-PSO-BBA	(5 nodes)	2.712	107	1.239	230
EDB-PSO-BBA	(5 nodes)	2.712	108	1.239	229
RR-PSO-BBA	(10 nodes)	2.712	104	1.239	202
MRR-PSO-BBA	(10 nodes)	2.712	93	1.239	206
DB-PSO-BBA	(10 nodes)	2.712	87	1.239	207
EDB-PSO-BBA	(10 nodes)	2.712	88	1.239	205

5.4.1 Adding more nodes to the equation

It is obvious that increasing the number of nodes makes the attack more evasive without an added efficiency cost. In the most extreme scenario, every query will be submitted using its own designated node. This ultimately causes no detections to occur. In fact, this scenario is a complete overkill. A detection is based on the average distance to the k nearest neighbors in the detector buffer associated with the account submitting the query. This means that the defense has to build up a history of k queries before any detection can occur. Only the $(k + 1)$ th and subsequent queries can therefore be flagged¹. Every adversarial attack that is distributed in the same manner, therefore only requires $\lceil budget/k \rceil$ nodes in order to remain undetected, where *budget* represents the query budget.

The combination of **PSO** and **BBA** as described in section 5.3 requires less nodes in order to be completely evasive. On the MNIST dataset, 25% of the adversarial attacks does not trigger any detections when distributed over 25 nodes. For 50 nodes, already 60% of the attacks are completely evasive and all attacks are completely evasive when using 100 nodes. The CIFAR-10 dataset requires more nodes to be

¹This limitation should not have a big impact on the defense against current state of the art methods, as **HSJA** is only 70% successful using 1000 queries [56] and the proposed value for k is 50 in [58]. In [58], being successful is defined as a maximum of 8/255 difference in L_∞ -distance between the original and the adversarial example.

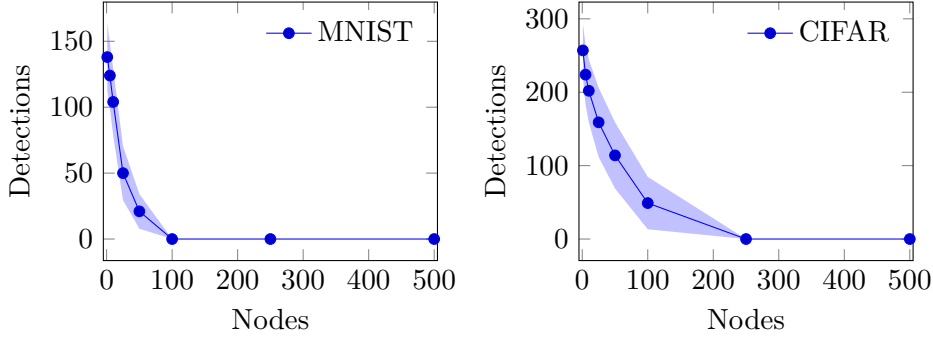


FIGURE 5.7: The number of detections for both datasets, when distributed over a given number of nodes. The number of detections decreases with increasing number of nodes. No attacks are flagged at 100 and 250 nodes for the MNIST and CIFAR dataset respectively. At 500 nodes, no detections can be flagged due to the theoretical limitation of the stateful defensive mechanism. The 95% confidence regions are shown around the curves. Note that the y axes do not line up.

completely evasive. When using 100 nodes, 60% of the attacks do not trigger any detections. If the node count is increased to 250 then no detections occur. Figure 5.7 shows the detection count in function of the number of nodes for both datasets. When the attack is distributed over 500 nodes, no detections are reported due to the limitations of the stateful defense mechanism.

5.4.2 Modifications to the distance based schemes

An improvement on the **DB** and **EDB** schemes can be made based on the assumption that the attacking user has knowledge about the status of its account. If the user receives a notification whenever its account is banned, then this knowledge can be incorporated in the distribution scheme. The standard **DB** and **EDB** schemes distribute queries based on the distance to the most recently submitted queries from this node. If an attack is flagged based on a submitted query, then this means that the corresponding account is banned and a new account has to be set up on the node. The new account will have a fresh detector buffer, therefore the previous queries submitted from this node are no longer relevant. By flushing the list of previously submitted queries whenever an account is banned, the distribution can happen more effectively. This distribution scheme will be called the **Flushing-Embedded-Distance-Based (FEDB)** distribution scheme. The results compared to the **EDB** scheme are shown in Table 5.6.

The **FEDB** scheme is more evasive than its non-flushing counterpart, both for five and ten nodes. However the gain is marginal. It is also based on the assumption that the attacker is notified of its account being banned. This might not always be the case. Users might be shadow banned [72] momentarily or even permanently, which

5. EVALUATION

TABLE 5.6: Results for the **FEDB** distribution scheme compared to **EDB** on the MNIST and CIFAR-10 dataset. The **FEDB** scheme outperforms its non-flushing counterpart by a slight margin.

Attack		MNIST		CIFAR	
		Distance	Detections	Distance	Detections
Baseline BBA		2.807	413	1.306	474
PSO-BBA		2.712	138	1.239	257
EDB-PSO-BBA	(5 nodes)	2.712	108	1.239	229
FEDB-PSO-BBA	(5 nodes)	2.712	105	1.239	224
EDB-PSO-BBA	(10 nodes)	2.712	88	1.239	205
FEDB-PSO-BBA	(10 nodes)	2.712	84	1.239	197

TABLE 5.7: Results for the **FEDB** distribution scheme for different buffer sizes on the MNIST and CIFAR-10 dataset. All experiments are performed using ten nodes. Increasing the buffer size has a slight advantage in terms of detections.

Attack		MNIST		CIFAR	
		Distance	Detections	Distance	Detections
Baseline BBA		2.807	413	1.306	474
PSO-BBA		2.712	138	1.239	257
FEDB-PSO-BBA	(buffer size 1)	2.712	91	1.239	208
FEDB-PSO-BBA	(buffer size 20)	2.712	84	1.239	197
FEDB-PSO-BBA	(buffer size 50)	2.712	82	1.239	194

violates the assumption².

The distance based distribution schemes can be greatly impacted by the length of the history they take into account. Longer buffers require more computational overhead, but the best node can be chosen with higher accuracy. Table 5.7 shows the results for different buffer lengths for distribution over ten nodes. Increasing the buffer, lowers the number of detections. However, the extra computations almost triple the run-time of the attack using a buffer size of 50 compared to a buffer of length 20. A buffer size of 20 is therefore an ideal middle ground between evasiveness and computational cost.

²A shadow ban detection algorithm should be implemented at the attack's side in case of a permanent shadow ban.

5.5 Throwing the defense off the scent

As previously mentioned, the two goals of the proposed attack are conflicting. In order for the attack to be efficient, subsequent queries will be clumped together in search space, causing the evasiveness of the attack to drop. This is an important trade-off that has to be made depending on the type of attack, cost of detection, efficiency thresholds and many other criteria. In this subsection changes will be made to the algorithm to be more evasive at the cost of efficiency.

Remember that the model under attack holds a query-bounded detector buffer per account. The buffer contains a history of the last submitted queries by a specific account. Every time a new query is inserted in the buffer, the distances to their nearest neighbors is determined and averaged. In the case of an attempted attack, the nearest neighbors will be relatively close together, causing the average distance to fall below a set threshold. Appropriate actions, such as banning the account, can now be taken.

The key idea behind the approach to increase the evasiveness is the following. Benign queries can be submitted by the attacker in order to flood the buffer with noise. The insertion of the benign queries will cause the average inter query distance to be raised, which in turn will lower the number of detections.

The approach raises two questions that will need to be answered: what sort of noise should be inserted and when should it be inserted. The former question will be answered first. Five types of noise are considered. The first one being uniform noise, as has been depicted in Figure 3.1a. The second and third options are Perlin noise with a fixed frequency value and a random frequency value respectively. Perlin noise has been shown in Figure 3.1b, while Figure 3.2 shows the influence of the frequency on the Perlin patterns. The fourth option will be random samples from the test data of the dataset on which the attack is being performed. The fifth and final option is a combination of the previous four strategies. Every time a benign query has to be inserted, one of the four options is chosen at random.

A similarity encoder will encode the submitted queries in an embedded space. The detector buffer will calculate the distances in this embedded space. Therefore it makes sense to determine the strategy with the highest inter query distance in the embedded space. The following experiment has been set up in order to establish the best strategy. The similarity encoder used in the EDB distribution scheme is reused to encode the benign queries. A fixed amount of queries (20 000) is encoded and added to a detector buffer of size 1000. Every time a query is added, the nearest neighbors are determined and the average distance to these neighbors is calculated. These distances are finally averaged over all submitted queries. The results of this experiment can be found in Table 5.8.

The most promising noise strategy is the test data strategy. This makes sense as the

TABLE 5.8: The average inter query distance (L_2 -distance in embedded space) for different noise strategies. Higher values correspond to better strategies as a detector is less likely to flag a potential attack.

	MNIST	CIFAR
Uniform noise	0.140	0.023
Perlin noise (fixed frequency)	0.110	0.004
Perlin noise	0.108	0.004
Test data	0.212	0.265
Mixed noise	0.238	0.198

similarity encoder is trained to maximize the distance between images representing different concepts. The uniform and perlin noise patterns are probably very similar to the encoder, although they might be further apart in terms of L_2 -distance. This causes these patterns to reside in the same regions of the embedded space. The mixed strategy is another promising candidate. It even outperforms the test data strategy on the MNIST dataset. However this is likely due to a large distance between the patterns and test data clusters in the embedded space as the mixed strategy triggers more detections (3056) compared to the test data strategy (0). For this reason, test data will be chosen as the source for the noise queries.

Now all that remains is answering the second question: when should the noise be inserted? Three strategies were considered. The first strategy is insertion of noise queries every n th query. The second strategy is inspired by the conclusion drawn from Figure 5.4. Detections are more common near the end of the attack, when the different particles have converged to the same region in search space. Therefore the second approach makes use of a geometric progression scheme [73] in order to submit more noise queries later in the attack progress. The third and final strategy does not have a fixed interval between the submission of noise queries. The submission is based on the distance to the previous queries. Using the EDB distribution scheme, a node is selected based on the average distance to the queries previously submitted from this node. If this average distance falls below a certain threshold then it is likely that an attack will be flagged by this query. Therefore noise queries are inserted before submitting the query used to progress the attack.

The first strategy does not seem a very good fit on paper. There is no good value for the submission interval n that works for the entire duration of the attack. Smaller values will result in noise being inserted too often, especially in the beginning of the attack when detections occur less frequently. Larger values will not help towards the evasiveness of the attack as the noise queries are inserted too infrequently. The second strategy relies on two hyperparameters, the starting value of the noise submission interval n and the decay rate. The optimal values of these parameters are very problem specific. For these reasons, the first two approaches are dropped from the evaluation.

TABLE 5.9: Results for the **EDB** distribution scheme with noise insertion for threshold values on the MNIST and CIFAR-10 dataset. All experiments are performed using ten nodes and a buffer size of 20.

Attack	Threshold	MNIST		CIFAR	
		Distance	Detections	Distance	Detections
Baseline BBA	/	2.807	413	1.306	474
PSO-BBA	/	2.712	138	1.239	257
EDB-PSO-BBA	/	2.712	88	1.239	205
EDB-PSO-BBA	0.10	2.710	87	1.537	208
EDB-PSO-BBA	0.25	2.719	87	1.622	202
EDB-PSO-BBA	1.00	2.720	83	1.595	221
EDB-PSO-BBA	2.00	2.775	89	1.592	241

Experiments will be performed using the third strategy. Noise queries will be inserted based on the mean distance of an incoming query to the previously submitted queries. Whenever this mean distance drops below a set threshold, a noise query will be inserted. The results of the experiments can be found in Table 5.9 for different threshold values. The experiments are performed using a **EDB** distribution scheme over ten nodes with a buffer size of 20. Lower values for the threshold cause less noise queries to be inserted, while higher value have the opposite effect. When only taking the number of detections into account, there seems to be an optimal dataset-dependent value. However, the noise insertion hampers the efficiency of the attack. Only one datapoint is more efficient than the algorithm without noise insertion due to random effects. All other datapoints are less efficient. For the CIFAR dataset, the attack with insertion performs worse than the vanilla **BBA** for all threshold values. The small gain in evasiveness is not worth the loss in efficiency, therefore no further experiments have been performed.

The queries inserted in the experiments performed in the previous paragraph are merely noise queries. They have no purpose in progressing the attack, only to evade the detection algorithm. What if the queries that are inserted to evade detection could also serve towards progressing the attack? The attacker would not have to worry about the query budget since all submitted queries would be useful. The idea behind this is based on the assumption that an adversary often interested in multiple adversarial examples instead of just the one. If this is the case, multiple attacks can be performed concurrently and their respective queries can be submitted in an interleaved manner.

The same experiments are performed again. This time experiments are grouped together. Experiments in the same group are executed concurrently and their queries that have be be submitted are interleaved. The query budget per experiment is

respected, meaning that a group of size s will submit s times the query budget. The interleaved queries are distributed over different nodes as if they came from a single attack.

Figure 5.8 shows the average number of detections per experiment. This means that the total number of detections of a group of size s is divided by s , which allows for an easier comparison to the baseline. Increasing the group size s yields a lower number of detections per experiment. The only exception to this conclusion is a group size of two for the CIFAR dataset. This effect is likely due to some unlucky combinations in the experiment pairings. Experiments with the same original label or same target label can cause more detections due to the attacks exploring the same regions in the search space. This effect is still present in the higher group sizes, but less prominent since the probability of all attacks exploring the same regions is reduced. The average L_2 -distance of the resulting adversarial examples is again equal to the results obtained by the **PSO-BBA** attack, since the only difference is the order in which the queries are submitted. All experiments performed to produce this plot used the **RR** distribution scheme over ten nodes.

In a real scenario, the attacker would benefit from interleaving the query submission of all planned adversarial attacks. There is an added benefit when the group size s is high enough. Since the detector buffer of the defensive scheme can only hold a certain number of queries, it should be possible that every individual attack only has a few of its queries at a time in this buffer. If this number is below the number of considered neighbors k , then no attack should be flagged even using a single node or account. This might not be possible with the MNIST or CIFAR datasets due to the limited number of target classes, causing the individual attacks to help detect each other. This effect can be seen on the CIFAR dataset with group size two in Figure 5.8. But when performing attacks on datasets such as ImageNet [57], where the number of classes is much higher, this scenario is not completely unthinkable.

5.6 Optimizing the attack

There are numerous parameters that influence both the efficiency and evasiveness of the attack. Some parameters are assigned values based on suggestions by the original authors of the work, some receive values based on experiments performed in this work and others get their respective value based on an educated guess. All parameters influence each other, making tuning of a single parameter pointless.

This section will focus on optimization of a small parameter set. Due to the complexity of this task, the goal is not to find the optimal value for all parameters in this set. But rather to find any obviously incompatible combinations of parameters if there are any. For the sake of completeness a list of all possible parameters and their respective default values is given in Table A.1 of appendix A. Table A.2 of the same




FIGURE 5.8: The average number of detections for both datasets, when experiments are combined and distributed over ten nodes. Experiments are run individually but their queries have to be submitted are interleaved. Increasing the number of experiments that are combined, decreases the average number of detections per experiment. The only data point inconsistent with this conclusion is two combined experiments for the CIFAR dataset. This is probably due to some unlucky combinations of experiments (same original label or same target label). Note that the y axes do not line up.

appendix shows the list of all parameters of the defensive scheme. The attacker has no knowledge of these values but the optimal values of the attacking parameters will depend on them.

5.7 Comparing to state of the art

As stated in section 3.2, **HSJA** [56] eclipsed **BA** and **BBA** in terms of efficiency. This section aims to compare **HSJA** to the attack proposed in this work both in terms of efficiency as evasiveness. The idea of query submission distribution is applicable to all adversarial attacks. In order to make the comparison more fair, this idea has been applied to **HSJA** as well.

Table 5.10 shows the results of the comparison. The original images and target labels are identical between both attacks and the results are averaged over all experiments. As expected, **HSJA** is more efficient than vanilla **BBA** and the spin-off **PSO-BBA** attack. For the CIFAR dataset, **PSO-BBA** performs more than four times worse when only taking the distance into account. **PSO-BBA** shows its potential when comparing the evasiveness of the attacks. This attack triggers less detections than **HSJA** both in the undistributed and distributed version. The increased evasiveness is partly due to the inherent distributed nature of the attack, but can also be attributed to the lower efficiency. The larger distance to the original image requires the attack to be less precise in order to remain adversarial. This in turn causes less detections. The larger distance to the original image does not have to pose a problem as the

5. EVALUATION

TABLE 5.10: Results for the **PSO-BBA** algorithm proposed in this work and the **HSJA** from [56]. **HSJA** is the more efficient attack, but triggers much more detections than **PSO-BBA**. Increasing the number of nodes causes a drop in the number of detection for both attacks.

Attack	Nodes	MNIST		CIFAR	
		Distance	Detections	Distance	Detections
Baseline BBA	1	2.807	413	1.306	474
PSO-BBA	1	2.712	138	1.239	257
HSJA	1	2.056	481	0.307	487
RR-PSO-BBA	10	2.712	104	1.239	202
HSJA	10	2.056	417	0.307	456

human eye is not very perceptible to such a small difference in L_2 -distance. This can be concluded from the visual comparison in Figures 5.9 and 5.10. The difference is notable for some examples, such as the airplane in Figure 5.10, but others it is not visible, i.e. the frog in the same figure.




FIGURE 5.9: The examples of Figure 5.1 are repeated with the corresponding adversarial examples. The first row are the original images with their respective labels above them. The second row is the result of the **PSO-BBA** attack with the target label specified above the resulting adversarial example. The L_2 -distance is shown after the target label in brackets. The setup of the third row is similar to the second row, but the adversarial examples originate from **HSJA**.




FIGURE 5.10: The examples of Figure 5.2 are repeated with the corresponding adversarial examples. The first row are the original images with their respective labels above them. The second row is the result of the **PSO-BBA** attack with the target label specified above the resulting adversarial example. The L_2 -distance is shown after the target label in brackets. The setup of the third row is similar to the second row, but the adversarial examples originate from **HSJA**.

Chapter 6

Conclusion

Appendices

Appendix A

Parameter table

This appendix serves as an overview of the different parameters present in the attack (and defense).

TABLE A.1: Parameter table

Parameter	Default value	Usage
Number of particles	5	Controls the number of particles in the PSO swarm. Increasing the number of particles decreases the evasiveness of the algorithm but decreases the efficiency as seen in section 5.3. Decreasing the value decreases both the evasiveness and efficiency.
Number of nodes N	10	Controls the number of nodes over which the query submission can be distributed. Increasing the number of nodes helps the algorithm to be more evasive as seen in section 5.4.1. Every added node comes with a certain cost.
Query budget	25000	Controls the maximum number of queries that can be submitted by the attack. Increasing the number can make the attack more efficient, but the later queries tend to trigger more detections, causing the evasiveness to drop.

A. PARAMETER TABLE

Parameter	Default value	Usage
Group size s	1	Controls the number of experiments that are grouped together and are executed concurrently. The queries of the attacks in the same group are submitted in an interleaved fashion. Increasing this value aids towards the evasiveness of the attack without sacrificing efficiency. Experiments have been performed with this parameter in section 5.5.
Distribution scheme	RR distribution	Controls how the queries are distributed over the different nodes. As discussed in section 5.4, the distribution scheme does not heavily influence the number of detections. Other options for this parameter are the MRR , DB and EDB distribution schemes.
History length	20	Controls the size of the query submission history that is taken into account in the DB and EDB distribution schemes. As seen in section 5.4.2, increasing this value helps the algorithm remain more evasive at an added computational cost. The chosen value is a middle ground between this computational cost and the evasiveness.
Source step ϵ	0.25 (MNIST) 0.20 (CIFAR)	Controls the step towards the original image in the BBA . This parameter is also shown in Figure 3.3 as ϵ . Increasing the value causes the queries to be more spread out over the search space, ultimately lowering detections. However, increasing the value too much will result in slow convergence of the attack.
Spherical step η	0.05	Controls the size of the random direction of the BBA algorithm. This parameter is also shown in Figure 3.3 as η . The value of this parameter is chosen based on the results of [54]. Increasing the value, adds more randomness into the algorithm, but can hamper convergence.

Parameter	Default value	Usage
Source step multiplier up	1.05	Controls the speed at which the source step increases when the new position remains adversarial. Setting this value to 1 (together with source step multiplier down) fixes the value of the source step for the entire duration of the attack. The increasing value of the source step helps the attack to be more evasive due to the submitted queries being more spread out.
Source step multiplier down	0.99	Controls the speed at which the source step decreases when the new position remains non-adversarial. Setting this value to 1 (together with source step multiplier up) fixes the value of the source step for the entire duration of the attack. The decreasing value of the source step helps the attack to be more efficient, since the smaller steps aid convergence.
Recalculate mask every	50	Controls after how many iterations the mask used in the BBA should be recalculated. This parameter has no big effect on the efficiency of the attack, only on the run time. After a sufficient amount of iterations the masks tends to change only slightly.
Particle acceleration coefficient c_p	2 ¹	Controls the rate of attraction towards the personal best position of a particle. The different steps of the particle are weighted according to equation 2.2. Increasing this value creates a stronger attraction to the personal best position of the particle.

¹ Due to the multi-group approach, these values are actually set based on equations 5.1 and 5.2 with $A1$ and $A2$ being 1 and 2 respectively.

A. PARAMETER TABLE

Parameter	Default value	Usage
Global acceleration coefficient c_g	1 ¹	Controls the rate of attraction towards the best position of the swarm. The different steps of the particle are weighted according to equation 2.2. Increasing this value creates a stronger attraction to the best position of the swarm.
Maximum velocity v_{max}	0.5	Controls the maximum velocity of a particle. The velocity values of equation 2.1 are clipped by this value in order to avoid the problem of exploding velocities. Decreasing this value causes the algorithm to take smaller steps in every iteration, which in turn reduces the evasiveness.
w_{start}	1	Controls the value of the inertia weight of equation 2.3. Decreasing the value will cause the particle to change its direction more easily in the beginning of the attack, allowing for more exploitation of the search space.
w_{end}	0	Controls the value of the inertia weight of equation 2.3. Increasing the value will cause the particle to change its direction more easily at the end of the attack, allowing for more exploitation of the search space.
Insert type	Test data	Controls the type of noise queries that are inserted in order to flush the buffer of the detector. The other options are uniforms noise, Perlin noise with a fixed frequency, Perlin noise with a variable frequency and a mix of all other approaches. Other approaches have a smaller inter query distance as shown in Table 5.8. Therefore train data noise is the best option in terms of evasiveness.

Parameter	Default value	Usage
Insert time	Distance based	Controls when noise queries should be inserted to flush the buffer of the detector. Other options are fixed interval and geometric decay.
Insert count	1	Controls the number of noise queries that are inserted every time the 'Insert time' decides that noise should be inserted. Inserting more noise queries can improve evasiveness, but the efficiency is lowered since less queries are used to progress the attack.
Insert threshold	/	Controls the threshold used in the distance based insertion scheme. Whenever the distance between an impending query and the previously submitted queries falls below this value, noise queries will be inserted. Increasing this value will result in more noise being inserted, causing the evasiveness to increase while the efficiency drops. Increasing the value too high will result in more detections due to the noise queries being flagged as an attack.
Insert decay rate	/	Controls the decay rate of the geometric decay insertion scheme. The number should be smaller than 1 in order to insert more queries near the end of the attack. The closer this number is to zero, the faster noise queries will be inserted.

A. PARAMETER TABLE

TABLE A.2: Parameter table of the defense²

Parameter	Default value	Usage
Buffer type	Query bounded	Controls how queries will be removed from the buffer. A query bounded buffer will hold a certain amount of queries. If a new query has to be added when the capacity is reached then queries are removed based on the first in, first out method [74]. The other option is a time bounded buffer where queries are removed once they reach a certain age. This approach requires more computations as checks need to happen regularly. Therefore the query bounded buffer is chosen in line with [58].
Buffer size	1000	Controls the number of queries the buffer holds before queries are removed again. Increasing the value results in more attack being possibly flagged but requires more storage. The value is chosen in line with [58].
k	50	Controls the number of neighbors that are considered when computing the average distance to a new incoming query. Increasing the value flags attacks with more certainty. However, this value is also equal to the number of queries that have to be inserted before any attack can be detected. Increasing the value too far will therefore result in no detections. The value is chosen in line with [58].

²Note that the attacker has no control over (and knowledge of) these parameters. However, the optimal parameters of the attack heavily depend on the defense parameters.

Parameter	Default value	Usage
Threshold	0.009 (MNIST) 0.021 (CIFAR)	Controls the threshold used to flag attacks. If the mean distance to the nearest neighbors of an incoming query falls below this value then the query is flagged as an attack in progress. Appropriate action can be taken against the account that submitted this query. The thresholds are determined based on the procedure discussed in [58]. A false positive rate of 0.1% will be achieved when adding all benign queries of the training set of the respective dataset.
Flush buffer after detection	True	Controls whether or not the buffer of the detector should be cleared after a detection happens. The value is set to true in order to simulate the user having to set up a new account after the previous account has been banned. If milder actions, such a temporary ban, are taken then the buffer should not be cleared.
Shadow ban	False	Controls whether or not the user is notified of a ban. If the value is false then the user is notified. This seems the most probable setting in real scenarios as there is no straightforward way to shadow ban a user while using an API. The attacker can use the knowledge of bans to its advantage as discussed in section 5.4.2.

Bibliography

- [1] F. Rosenblatt. The perceptron: A probabilistic model for information storage and organization in the brain. 65(6):386–408.
- [2] Vinod Nair and Geoffrey E. Hinton. Rectified linear units improve restricted boltzmann machines. In *Proceedings of the 27th International Conference on International Conference on Machine Learning, ICML'10*, page 807–814, Madison, WI, USA, 2010. Omnipress.
- [3] Wikipedia. Heaviside step function. https://en.wikipedia.org/wiki/Heaviside_step_function, 04 2022. Accessed: 2022-04-30.
- [4] Wikipedia. Softmax function. https://en.wikipedia.org/wiki/Softmax_function, 03 2022. Accessed: 2022-04-30.
- [5] M.V. Valueva, N.N. Nagornov, P.A. Lyakhov, G.V. Valuev, and N.I. Chervyakov. Application of the residue number system to reduce hardware costs of the convolutional neural network implementation. *Mathematics and Computers in Simulation*, 177:232–243, 2020.
- [6] Steve Lawrence, C. Lee Giles, Ah Chung Tsoi, and Andrew D. Back. Face recognition: A convolutional neural network approach. *IEEE TRANSACTIONS ON NEURAL NETWORKS*, 8(1):98–113, 1997.
- [7] Xiangang Li and Xihong Wu. Constructing long short-term memory based deep recurrent neural networks for large vocabulary speech recognition. *CoRR*, abs/1410.4281, 2014.
- [8] Ilya Sutskever, Oriol Vinyals, and Quoc V. Le. Sequence to sequence learning with neural networks. *CoRR*, abs/1409.3215, 2014.
- [9] Hansika Hewamalage, Christoph Bergmeir, and Kasun Bandara. Recurrent neural networks for time series forecasting: Current status and future directions. *CoRR*, abs/1909.00590, 2019.
- [10] Jon Russell. Google’s alphago ai wins three-match series against the world’s best go player. <https://techcrunch.com/2017/05/24/alphago-beats-planets-best-human-go-player-ke-jie/amp/?guccounter=1>, 05 2017. Accessed: 2021-12-08.

- [11] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems*, volume 25. Curran Associates, Inc., 2012.
- [12] Wikipedia. Perceptron. <https://en.wikipedia.org/wiki/Perceptron>, 2022 03. Accessed: 2022-04-16.
- [13] Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [14] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. Intriguing properties of neural networks, 2014.
- [15] Fatemeh Vakhshiteh, Ahmad Nickabadi, and Raghavendra Ramachandra. Adversarial attacks against face recognition: A comprehensive study, 2021.
- [16] Abhiram Gnanasambandam, Alex M. Sherman, and Stanley H. Chan. Optical adversarial attack, 2021.
- [17] Octavian Suciu, Scott E. Coull, and Jeffrey Johns. Exploring adversarial examples in malware detection, 2019.
- [18] Nicholas Carlini and David Wagner. Audio adversarial examples: Targeted attacks on speech-to-text, 2018.
- [19] Giulio Zizzo, Chris Hankin, Sergio Maffei, and Kevin Jones. Invited: Adversarial machine learning beyond the image domain. In *2019 56th ACM/IEEE Design Automation Conference (DAC)*, pages 1–4, 2019.
- [20] Victoria Drake. Threat modeling. https://owasp.org/www-community/Threat_Modeling#:~:text=A%20threat%20model%20is%20a,through%20the%20lens%20of%20security., 09 2021. Accessed: 2022-03-31.
- [21] Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples, 2015.
- [22] Nicholas Carlini and David Wagner. Towards evaluating the robustness of neural networks, 2017.
- [23] Florian Tramèr, Alexey Kurakin, Nicolas Papernot, Ian Goodfellow, Dan Boneh, and Patrick McDaniel. Ensemble adversarial training: Attacks and defenses. In *International Conference on Learning Representations*, 2018.
- [24] Wikipedia. Distance. <https://en.wikipedia.org/wiki/Distance>, 11 2021. Accessed: 2021-12-09.
- [25] Wikipedia. Discrete space. https://en.wikipedia.org/wiki/Discrete_space, 04 2022. Accessed: 2022-04-30.

-
- [26] Gabriel Resende Machado, Eugênio Silva, and Ronaldo Ribeiro Goldschmidt. Adversarial machine learning in image classification: A survey toward the defender’s perspective. *ACM Comput. Surv.*, 55(1), nov 2021.
 - [27] Gabriel Resende Machado, Eugênio Silva, and Ronaldo Ribeiro Goldschmidt. Adversarial machine learning in image classification: A survey towards the defender’s perspective. *CoRR*, abs/2009.03728, 2020.
 - [28] Nicolas Papernot, Patrick McDaniel, Ian Goodfellow, Somesh Jha, Z. Berkay Celik, and Ananthram Swami. Practical black-box attacks against machine learning. In *Proceedings of the 2017 ACM on Asia Conference on Computer and Communications Security*, ASIA CCS ’17, page 506–519, New York, NY, USA, 2017. Association for Computing Machinery.
 - [29] Anish Athalye, Nicholas Carlini, and David A. Wagner. Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples. *CoRR*, abs/1802.00420, 2018.
 - [30] Nicolas Papernot, Patrick D. McDaniel, Xi Wu, Somesh Jha, and Ananthram Swami. Distillation as a defense to adversarial perturbations against deep neural networks. *CoRR*, abs/1511.04508, 2015.
 - [31] Shixiang Gu and Luca Rigazio. Towards deep neural network architectures robust to adversarial examples, 2015.
 - [32] Dan Hendrycks and Kevin Gimpel. Visible progress on adversarial images and a new saliency map. *CoRR*, abs/1608.00530, 2016.
 - [33] Chuan Guo, Mayank Rana, Moustapha Cissé, and Laurens van der Maaten. Countering adversarial images using input transformations. *CoRR*, abs/1711.00117, 2017.
 - [34] Nicolas Papernot and Patrick D. McDaniel. Deep k-nearest neighbors: Towards confident, interpretable and robust deep learning. *CoRR*, abs/1803.04765, 2018.
 - [35] Xiaoyu Cao and Neil Zhenqiang Gong. Mitigating evasion attacks to deep neural networks via region-based classification. *CoRR*, abs/1709.05583, 2017.
 - [36] James Kennedy and Russell Eberhart. Particle swarm optimization. In *Proceedings of ICNN’95-international conference on neural networks*, volume 4, pages 1942–1948. IEEE, 1995.
 - [37] Marco Dorigo, Vittorio Maniezzo, and Alberto Coloni. Ant system: Optimization by a colony of cooperating agents. *IEEE transactions on systems, man, and cybernetics - part b. IEEE transactions on systems, man, and cybernetics. Part B, Cybernetics : a publication of the IEEE Systems, Man, and Cybernetics Society*, 26:29–41, 02 1996.

- [38] John H. Holland. *Genetic Algorithms and Adaptation*, pages 317–333. Springer US, Boston, MA, 1984.
- [39] Rami Abousleiman and Osamah Rawashdeh. Electric vehicle modelling and energy-efficient routing using particle swarm optimisation. *IET Intelligent Transport Systems*, 10(2):65–72, 2016.
- [40] Tadashi Yamada and Zukhruf Febri. Freight transport network design using particle swarm optimisation in supply chain–transport supernetwork equilibrium. *Transportation Research Part E: Logistics and Transportation Review*, 75:164–187, 2015.
- [41] Bruno Almeida and Victor Coppo Leite. *Particle Swarm Optimization: A Powerful Technique for Solving Engineering Problems*. 12 2019.
- [42] Romeela Mohee and Ackmez Mudhoo. Analysis of the physical properties of an in-vessel composting matrix. *Powder Technology*, 155(1):92–99, 2005.
- [43] Y. Shi and R.C. Eberhart. Empirical study of particle swarm optimization. In *Proceedings of the 1999 Congress on Evolutionary Computation-CEC99 (Cat. No. 99TH8406)*, volume 3, pages 1945–1950 Vol. 3, 1999.
- [44] M. Clerc and J. Kennedy. The particle swarm - explosion, stability, and convergence in a multidimensional complex space. *IEEE Transactions on Evolutionary Computation*, 6(1):58–73, 2002.
- [45] Wikipedia. Vanilla software. https://en.wikipedia.org/wiki/Vanilla_software, 04 2022. Accessed: 2022-05-12.
- [46] Naufal Suryanto, Chihiro Ikuta, and Dadet Pramadihanto. Multi-group particle swarm optimization with random redistribution. In *2017 International Electronics Symposium on Knowledge Creation and Intelligent Computing (IES-KCIC)*, pages 1–5, 2017.
- [47] Rainer Storn and Kenneth Price. Differential evolution: A simple and efficient adaptive scheme for global optimization over continuous spaces. *Journal of Global Optimization*, 23, 01 1995.
- [48] Radha Thangaraj, Millie Pant, Ajith Abraham, and Pascal Bouvry. Particle swarm optimization: Hybridization perspectives and experimental illustrations. *Applied Mathematics and Computation*, 217(12):5208–5226, 2011.
- [49] Eid Albalawi, Fujie Chen, Ruppa K. Thulasiram, and Parimala Thulasiraman. Effect of pso communication topologies on task matching in grid computing. *2019 IEEE Congress on Evolutionary Computation (CEC)*, pages 426–433, 2019.
- [50] Takahiro Tsujimoto, Takuya Shindo, Takayuki Kimura, and Kenya Jin’no. A relationship between network topology and search performance of pso. In *2012 IEEE Congress on Evolutionary Computation*, pages 1–6, 2012.

-
- [51] Wieland Brendel, Jonas Rauber, and Matthias Bethge. Decision-Based Adversarial Attacks: Reliable Attacks Against Black-Box Machine Learning Models. *arXiv:1712.04248 [cs, stat]*, February 2018. arXiv: 1712.04248.
- [52] Ye Wenhe. Trust-region methods. https://optimization.mccormick.northwestern.edu/index.php/Trust-region_methods, 06 2014. Accessed: 2021-12-09.
- [53] Alhussein Fawzi, Seyed-Mohsen Moosavi-Dezfooli, and Pascal Frossard. Robustness of classifiers: from adversarial to random noise. *CoRR*, abs/1608.08967, 2016.
- [54] Thomas Brunner, Frederik Diehl, Michael Truong Le, and Alois Knoll. Guessing Smart: Biased Sampling for Efficient Black-Box Adversarial Attacks. *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4957–4965, October 2019. arXiv: 1812.09803.
- [55] Ken Perlin. An image synthesizer. In *Proceedings of the 12th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH ’85, page 287–296, New York, NY, USA, 1985. Association for Computing Machinery.
- [56] Jianbo Chen, Michael I. Jordan, and Martin J. Wainwright. HopSkipJumpAttack: A Query-Efficient Decision-Based Attack. *arXiv:1904.02144 [cs, math, stat]*, April 2020. arXiv: 1904.02144.
- [57] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- [58] Steven Chen, Nicholas Carlini, and David Wagner. Stateful Detection of Black-Box Adversarial Attacks. *arXiv:1907.05587 [cs]*, July 2019. arXiv: 1907.05587.
- [59] Sean Bell and Kavita Bala. Learning visual similarity for product design with convolutional neural networks. *ACM Trans. Graph.*, 34(4), jul 2015.
- [60] Alex Krizhevsky. Learning multiple layers of features from tiny images. pages 32–33, 2009.
- [61] Moustafa Alzantot, Yash Sharma, Supriyo Chakraborty, Huan Zhang, Chojui Hsieh, and Mani Srivastava. GenAttack: Practical Black-box Attacks with Gradient-Free Optimization. *arXiv:1805.11090 [cs]*, June 2019. arXiv: 1805.11090.
- [62] Yinpeng Dong, Hang Su, Baoyuan Wu, Zhifeng Li, Wei Liu, Tong Zhang, and Jun Zhu. Efficient decision-based black-box adversarial attacks on face recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7714–7722, 2019.

- [63] Rayan Mosli, Matthew Wright, Bo Yuan, and Yin Pan. They might not be giants: crafting black-box adversarial examples with fewer queries using particle swarm optimization. *arXiv preprint arXiv:1909.07490*, 2019.
- [64] Hyunjun Mun, Sunggwan Seo, Baehoon Son, and Joobeom Yun. Black-box audio adversarial attack using particle swarm optimization. *IEEE Access*, 10:23532–23544, 2022.
- [65] Fengtao Xiang, Jiahui Xu, Wanpeng Zhang, and Weidong Wang. A distributed biased boundary attack method in black-box attack. *Applied Sciences*, 11(21), 2021.
- [66] Naufal Suryanto, Hyoeun Kang, Yongsu Kim, Youngyeo Yun, Harashta Tatimma Larasati, and Howon Kim. A distributed black-box adversarial attack based on multi-group particle swarm optimization. *Sensors*, 20(24), 2020.
- [67] Y. Shi and R.C. Eberhart. Empirical study of particle swarm optimization. In *Proceedings of the 1999 Congress on Evolutionary Computation-CEC99 (Cat. No. 99TH8406)*, volume 3, pages 1945–1950 Vol. 3, 1999.
- [68] Kaspersky. What is a ddos attack? - ddos meaning. www.usa.kaspersky.com/resource-center/threats/ddos-attacks, 01 2021. Accessed: 2022-03-18.
- [69] Google. Cloud Vision pricing. <https://cloud.google.com/vision/pricing#prices>, 04 2022. Accessed: 2022-04-08.
- [70] Amazon. Amazon Rekognition pricing. <https://aws.amazon.com/rekognition/pricing/>, 04 2022. Accessed: 2022-04-08.
- [71] Wikipedia. Round-robin scheduling. https://en.wikipedia.org/wiki/Round-robin_scheduling, 04 2022. Accessed: 2022-04-18.
- [72] Wikipedia. Shadow banning. https://en.wikipedia.org/wiki/Shadow_banning, 04 2022. Accessed: 2022-05-11.
- [73] Wikipedia. Geometric progression. https://en.wikipedia.org/wiki/Geometric_progression, 03 2022. Accessed: 2022-05-15.
- [74] Wikipedia. FIFO (computing and electronics). [https://en.wikipedia.org/wiki/FIFO_\(computing_and_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics)), 05 2022. Accessed: 2022-05-18.